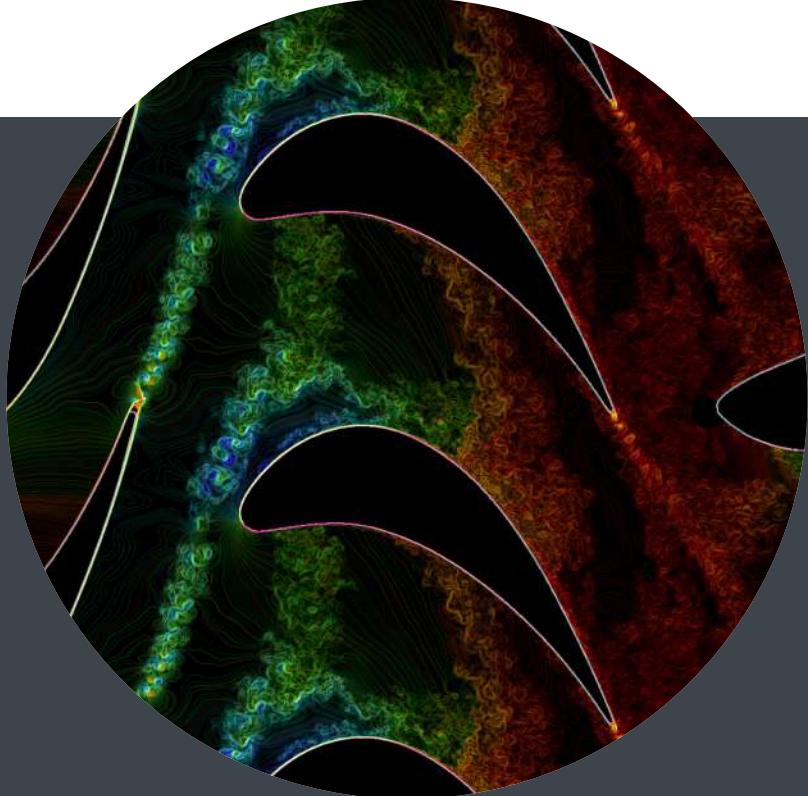


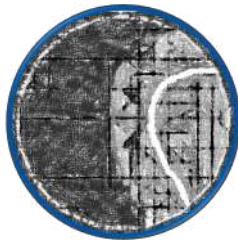
# Data-Driven, Discretization Consistent LES models

Andrea Beck

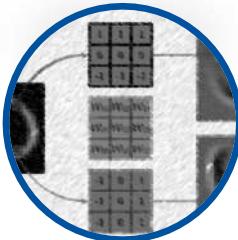
Institute of Aerodynamics and Gas Dynamics (IAG) &  
Stuttgart Center for Simulation Science (SC SimTech)



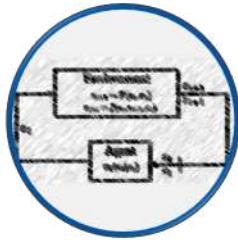
# Overview



Part 1: Multi-X problems and PDE solvers



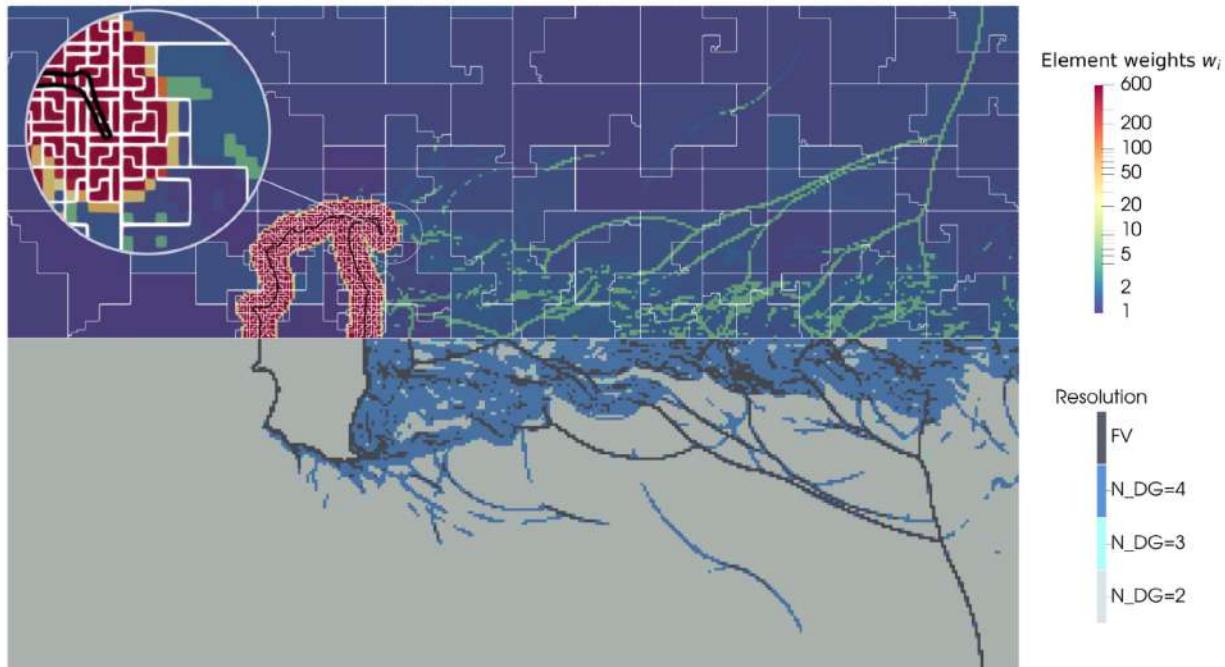
Part 2: Supervised learning for PDE submodels



Part 3: From data-driven to integrated CFD/ML

# Multiphase, Multiscale, Multimethod: Shock Droplet Interaction

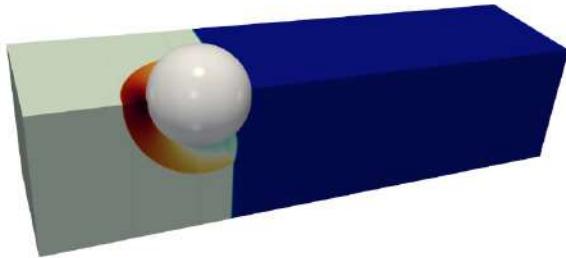
- Sharp-Interface simulation of 3D water droplet – shock interaction at  $Ma = 2.4$  and  $We = 100$
- About 140 DOF / droplet diameter
- hp-adaptive DGSEM / FV scheme with DLB
- 7200 CPUh on HAWK



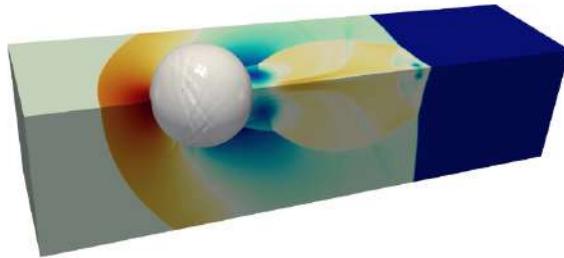
Mossier et al., An Efficient hp-Adaptive Strategy for a Level-Set Ghost Fluid Method, arxiv, 2023

Jöns et al., Riemann solvers for phase transition in a compressible sharp-interface method, JCP, 2023

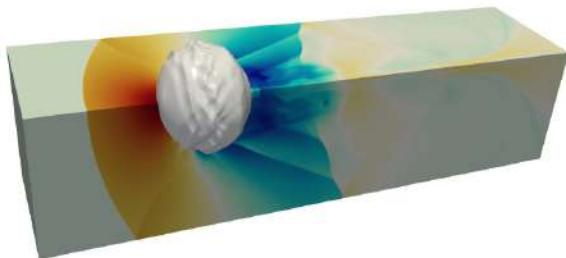
Mossier et al., A p-Adaptive Discontinuous Galerkin Method with hp-Shock Capturing. Journal of Scientific Computing, 2022



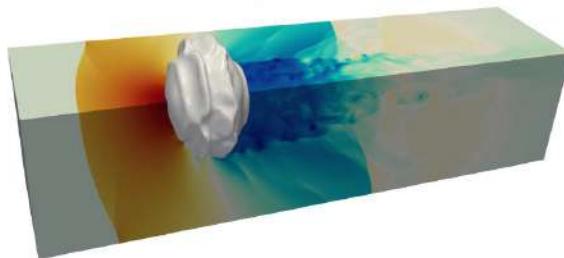
(a)  $t^* = 0.8$



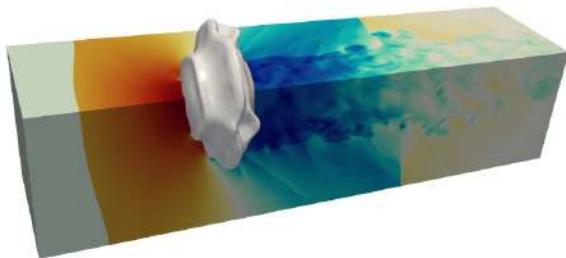
(b)  $t^* = 2.4$



(c)  $t^* = 4.7$



(d)  $t^* = 7.0$



(e)  $t^* = 9.4$



(f)  $t^* = 11.8$

# Multiphase, Multiscale, Multimethod: Icing on Airfoils



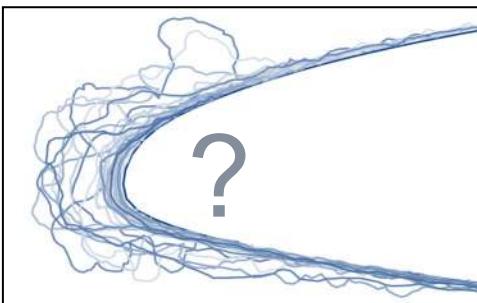
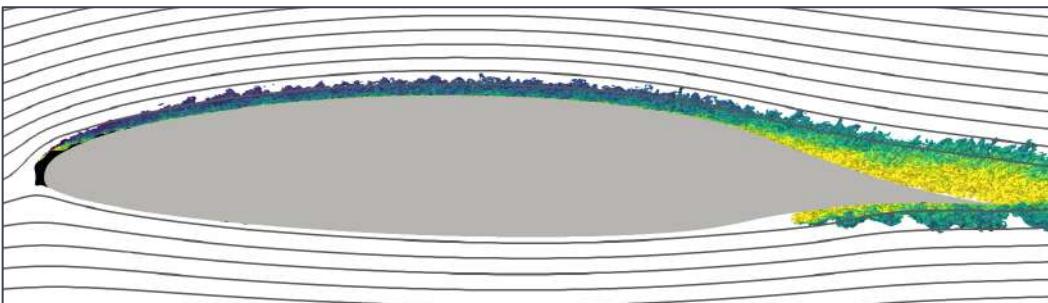
EIROS project via <http://www.eirosproject.com/project/case-study-applications/>



Stefano Fornasier via [http://tesi.cab.unipd.it/50432/1/STEFANO\\_FORNASIER\\_1147354\\_assignsubmission\\_file\\_Fornasier\\_Stefano\\_1056548.pdf](http://tesi.cab.unipd.it/50432/1/STEFANO_FORNASIER_1147354_assignsubmission_file_Fornasier_Stefano_1056548.pdf)

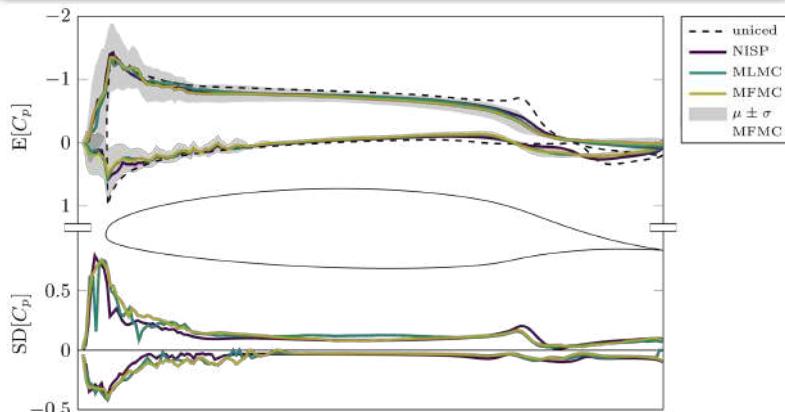
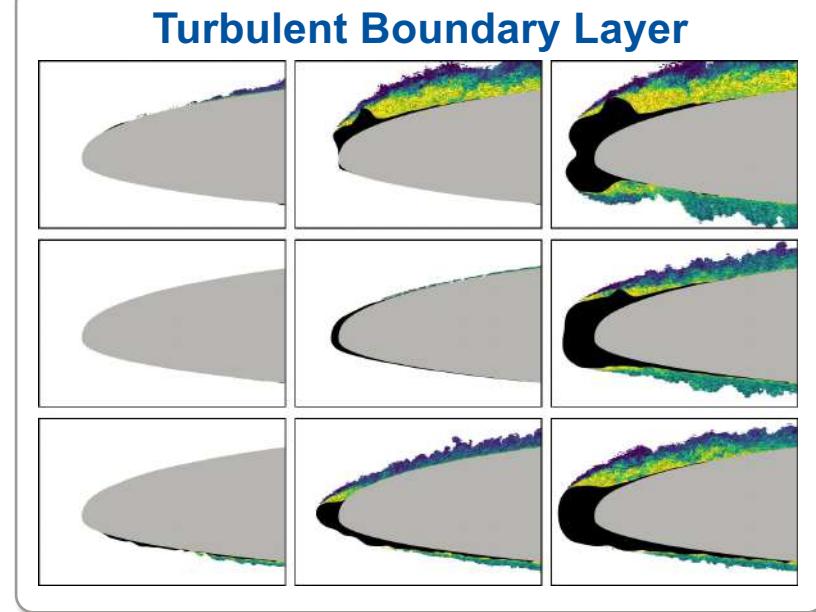


Chicago Tribune via <https://www.newspapers.com/clip/14788830/american-eagle-flight-4184-crashes-in-in/>

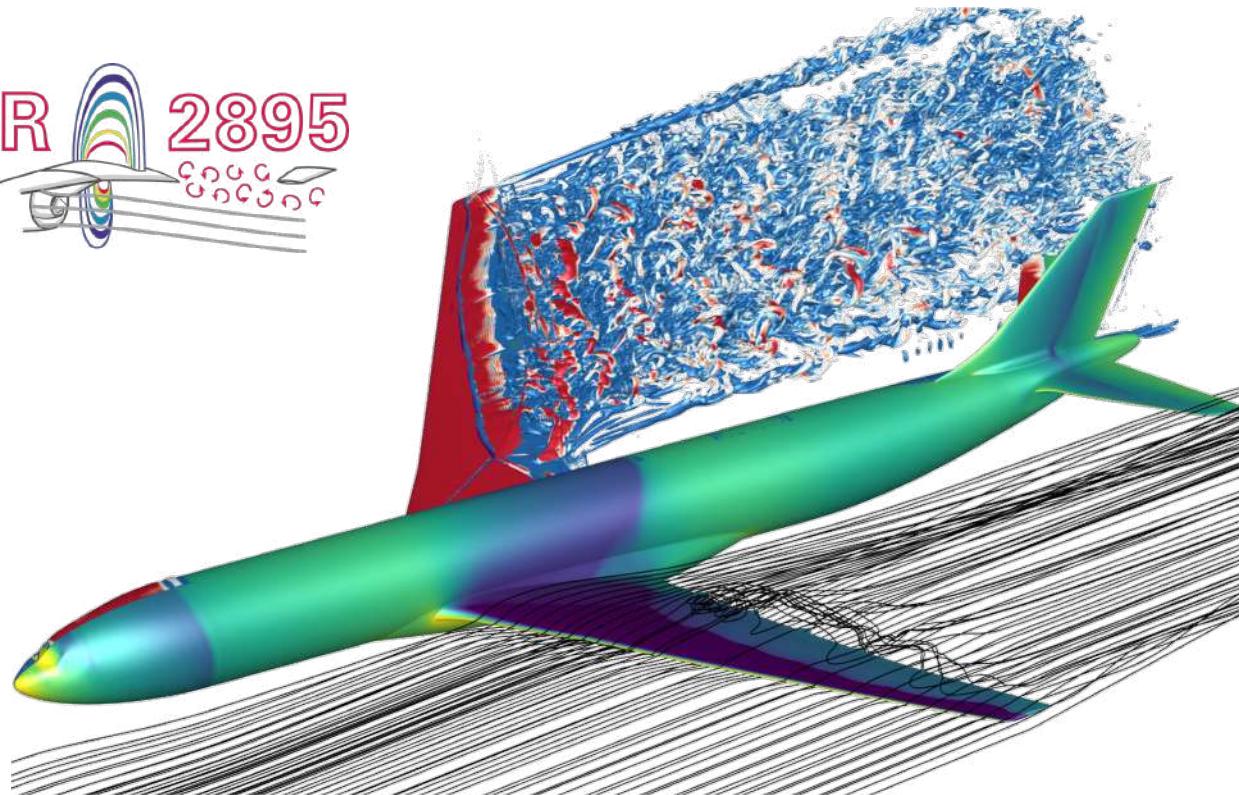


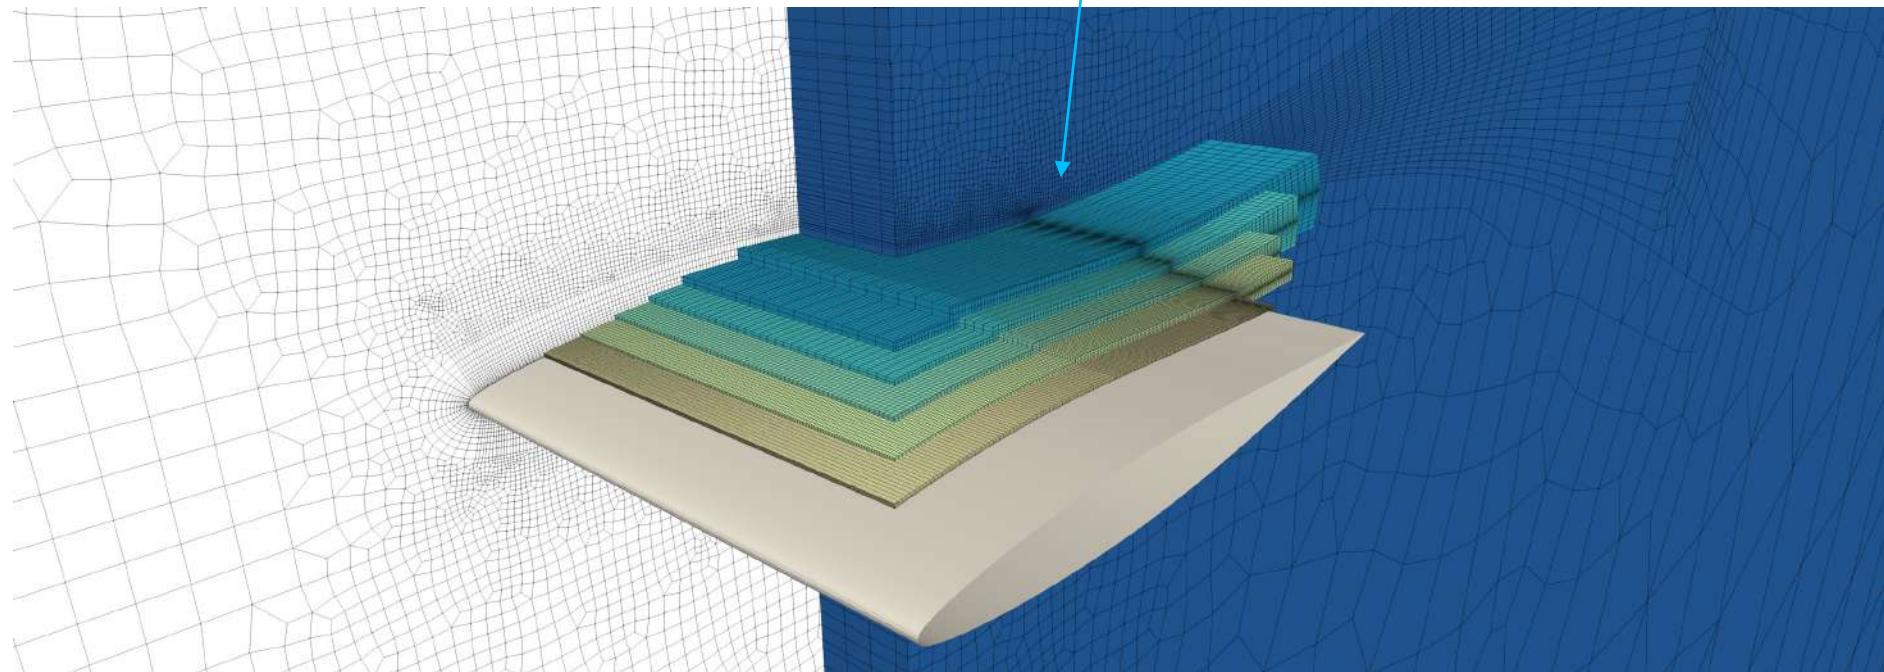
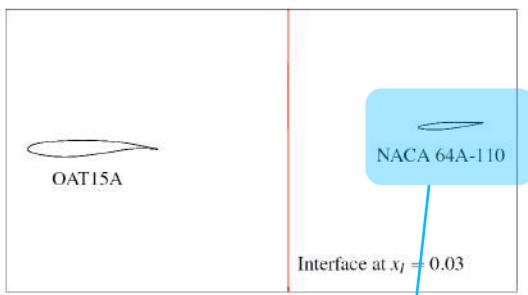
# Multiphase, Multiscale, Multimethod

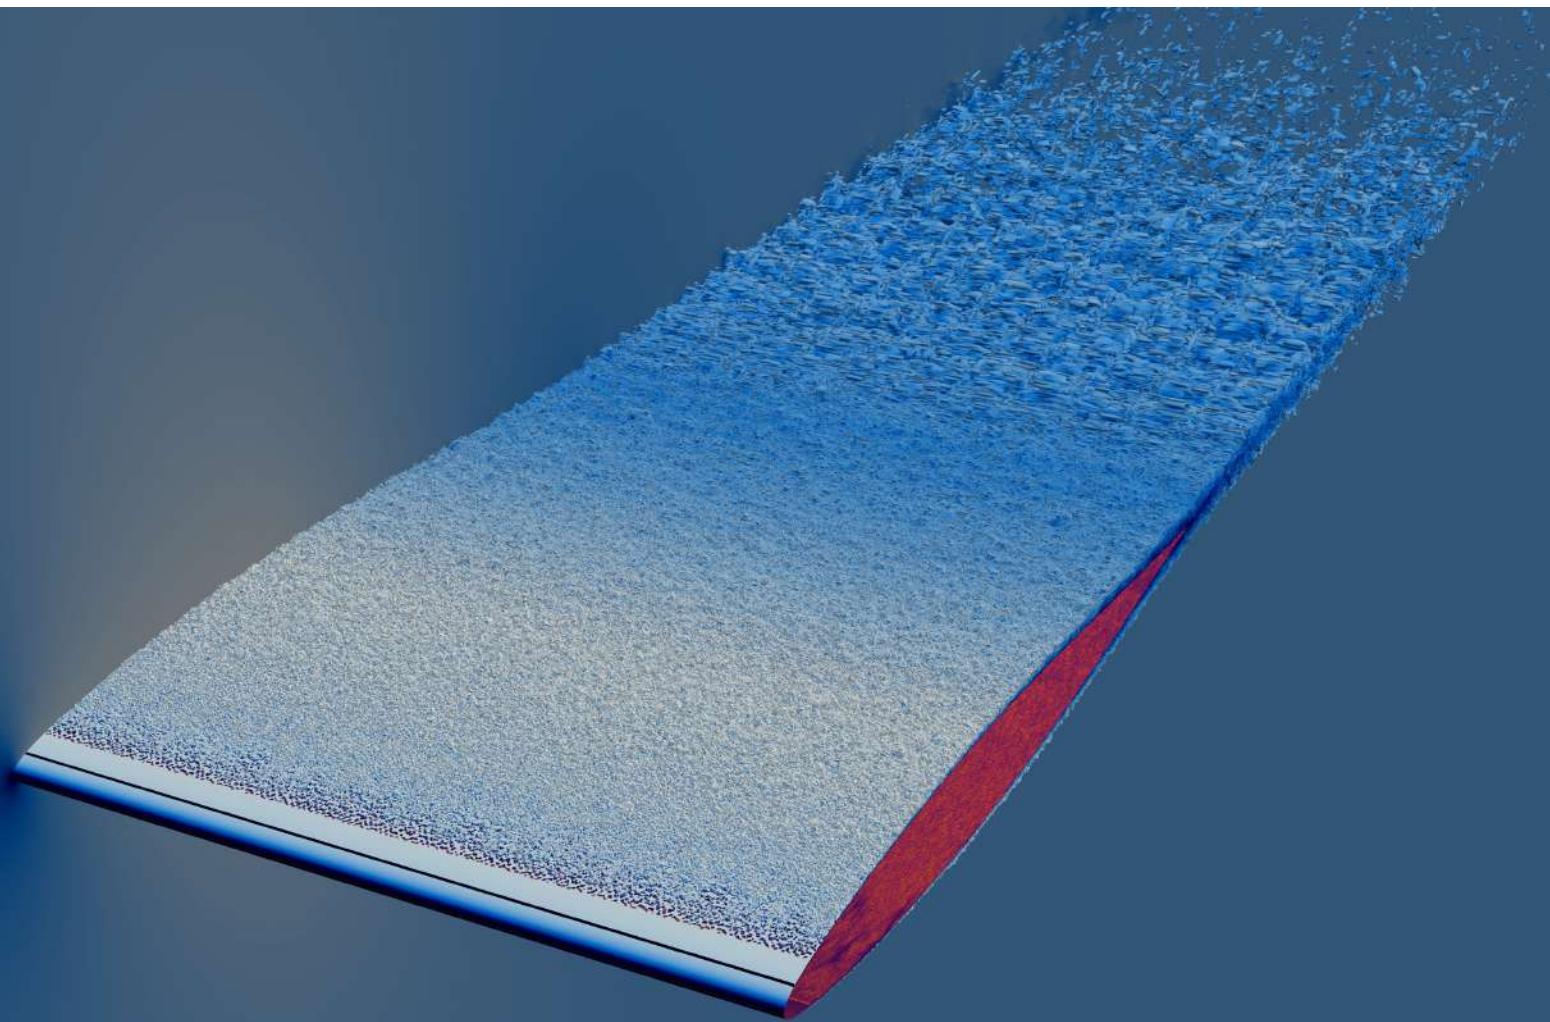
- Uncertainty Quantification Analysis of Icing on Airfoil Performance
- Parameterization of iced shapes via POD / PCA
- Automated grid generation for wall-resolved LES on hexa-only grids
- Non-intrusive Spectral Projection, MLMC and MFMC of 3D LES computations
- PoUnce-Framework for HPC scheduling and management of UQ workflow



# Multiphase, Multiscale, Multimethod: Wake / HTP interaction under buffet

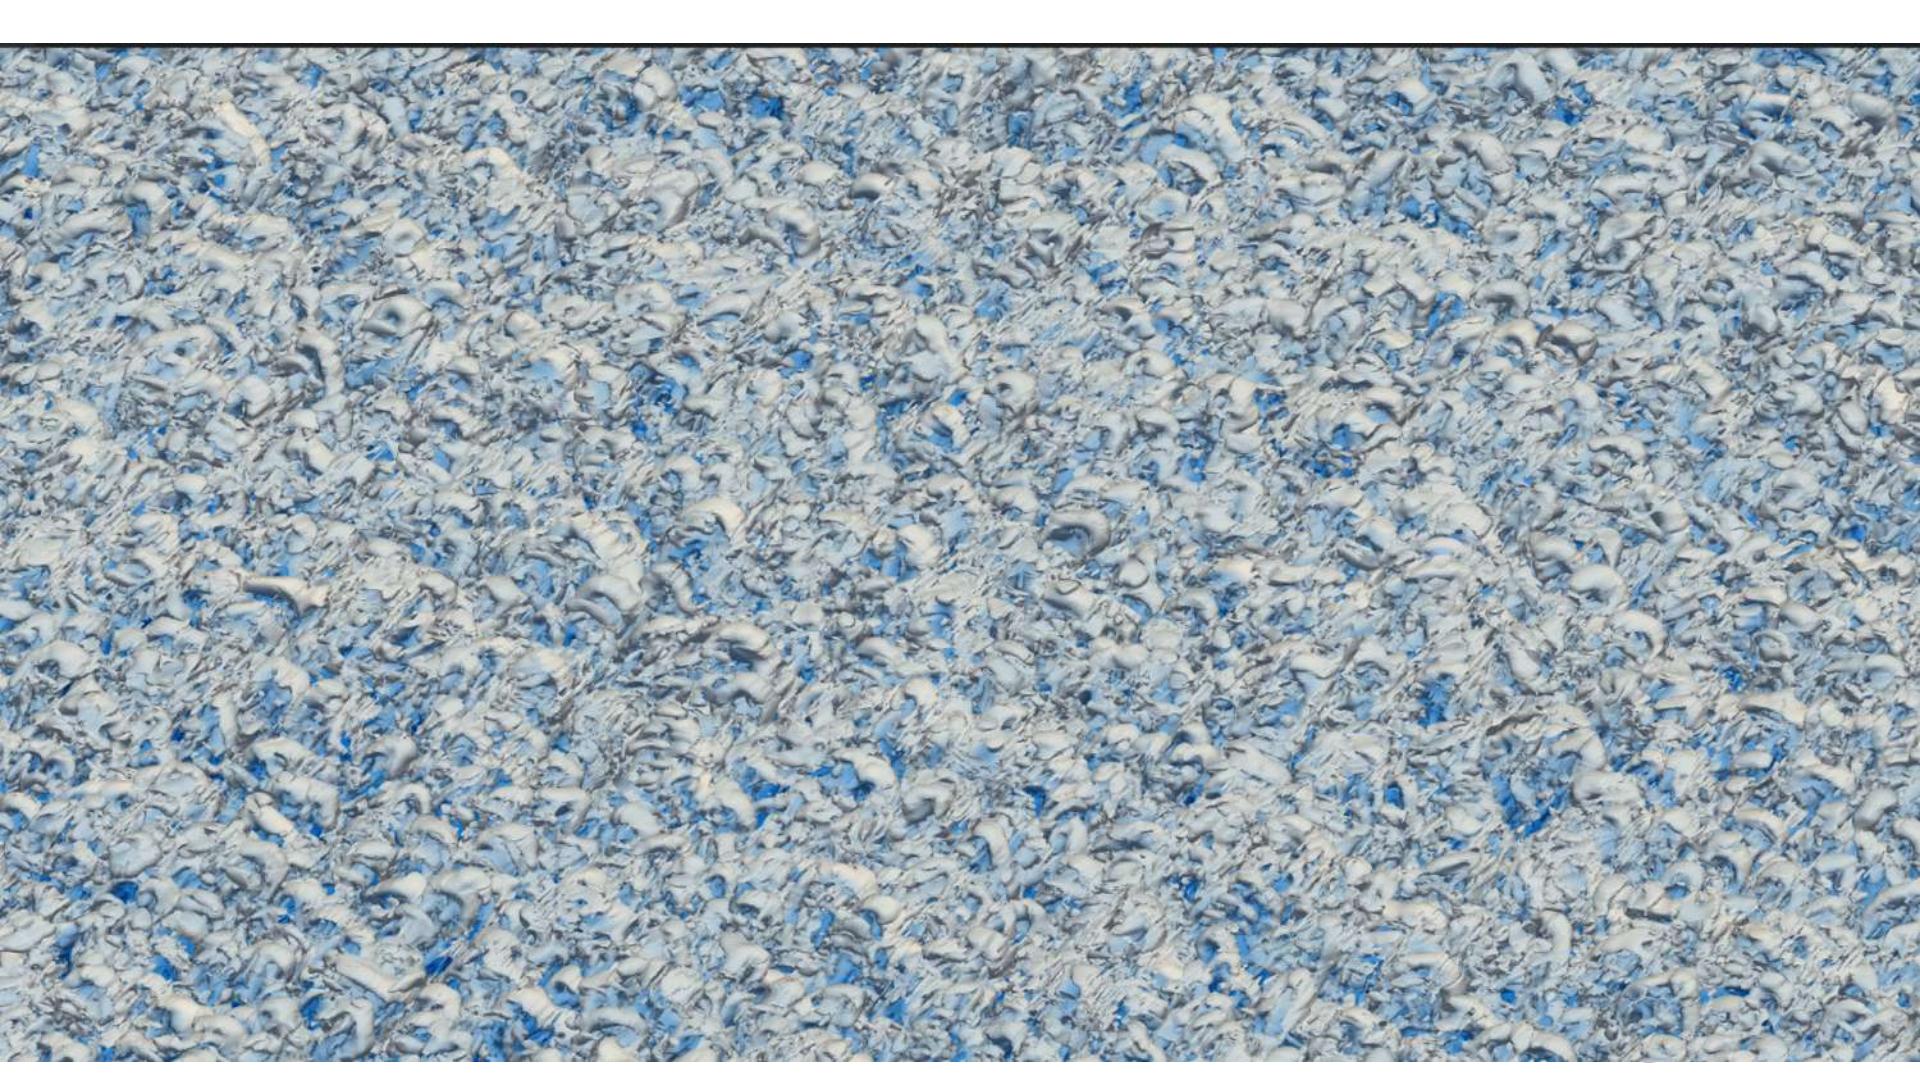


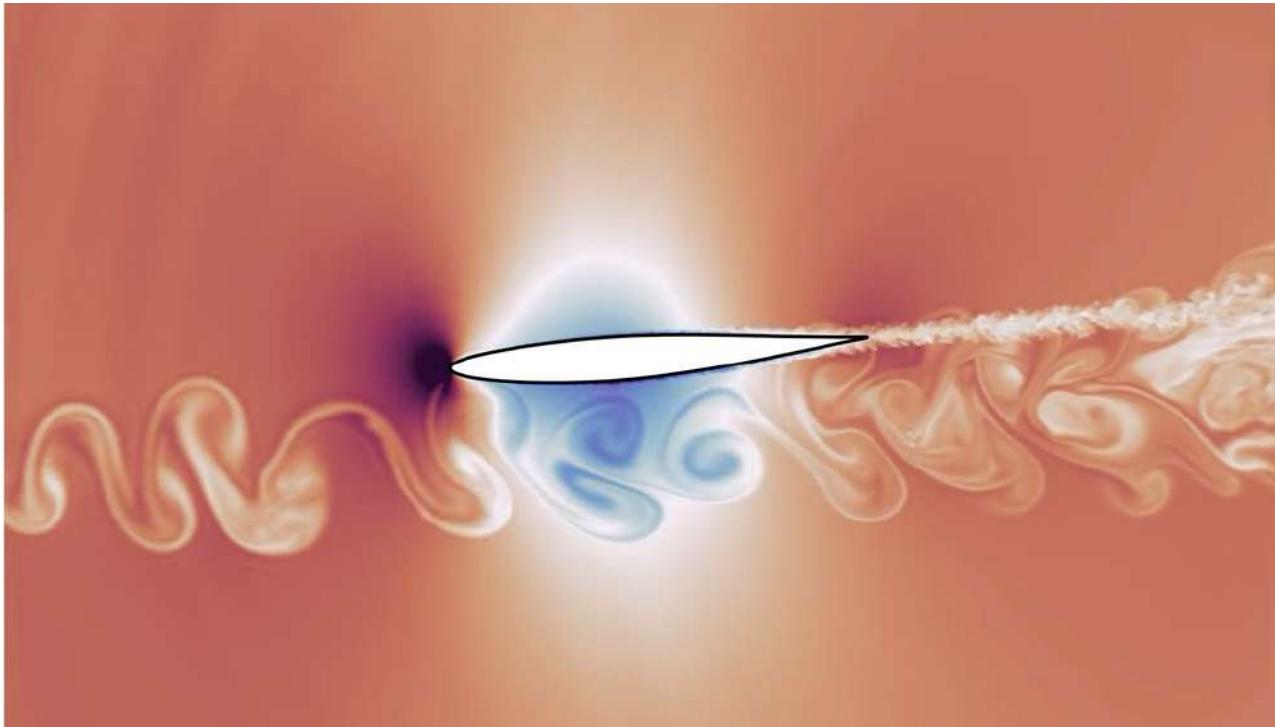




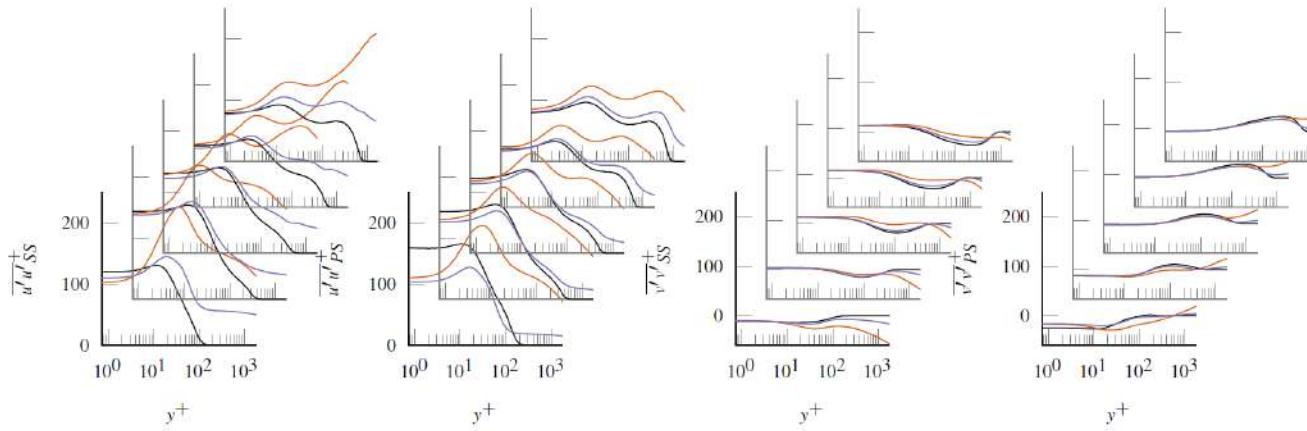






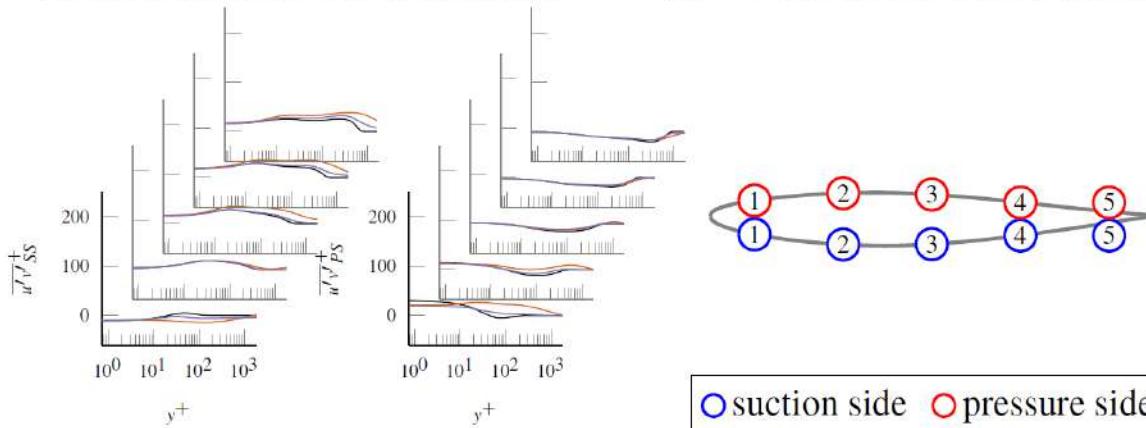


- Approx. 1 Billion DOF, N=7, 70 CTU



(c) Streamwise mean velocity fluctuations

(d) Wall-normal mean velocity fluctuations.



(e)  $\overline{u'v'^+}$  velocity fluctuation.

— steady — unsteady: Phase 1 — unsteady: Phase 2



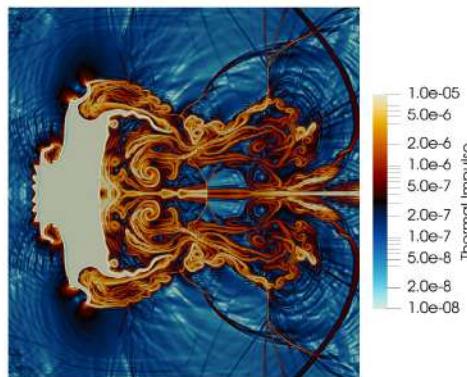
# What we need to solve these problems – without proof

- High order accurate schemes for sufficiently smooth regions: 2<sup>nd</sup> order FV will not do
- Monotone schemes for shocks, interfaces, gradients
- h/p/r/x-adaptivity to the underlying physics: A flexible, intelligent, physics-compatible numerical scheme

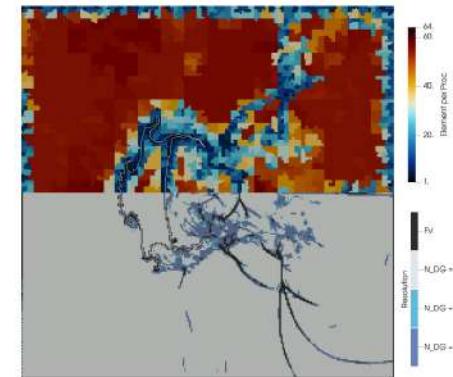


FLEXI: A high order discontinuous Galerkin framework for hyperbolic–parabolic conservation laws

<https://github.com/flexi-framework/flexi>

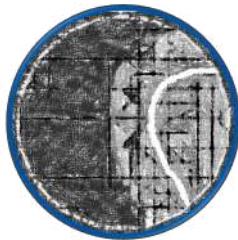


(a) Thermal impulse  $|\rho j|$

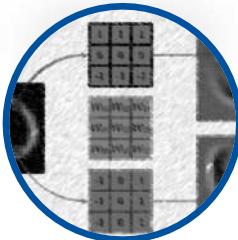


(b) Adaptive discretization

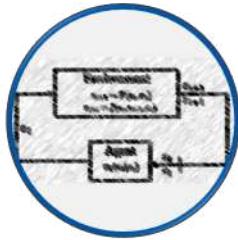
# Overview



Part 1: Multi-X problems and PDE solvers



Part 2: Supervised learning for PDE submodels



Part 3: From data-driven to integrated CFD/ML

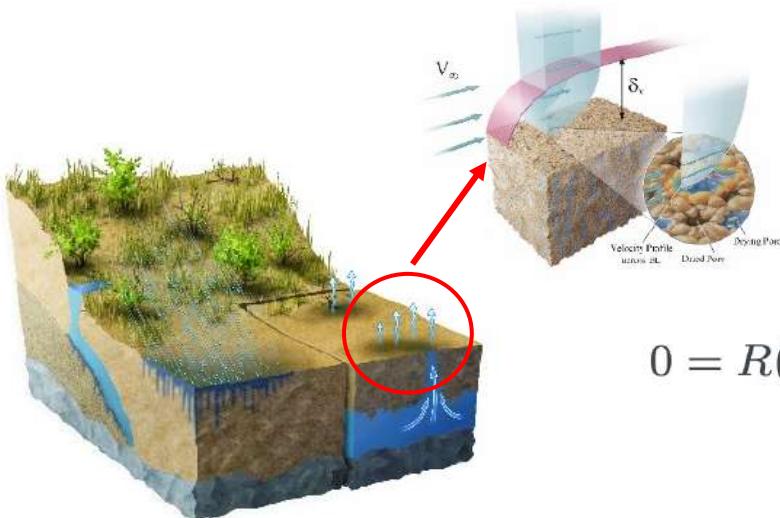
# The Scale Gap: Where and why do we need ML?

- Multiscale-Multiphysics systems can cover a **wide range of scales**



# The Scale Gap

- Multiscale problems: Formulations on micro, meso, macro, ... scales
- Definition of a **coarsening operator** (1) → (2) possibly ambiguous
- Solving coarse-scale equations plus Models for the **Closure terms  $M(U_h)$**



$$\underbrace{R(U_l, x_l) = 0}_{\text{Governing Equations at level l}} \quad (1)$$

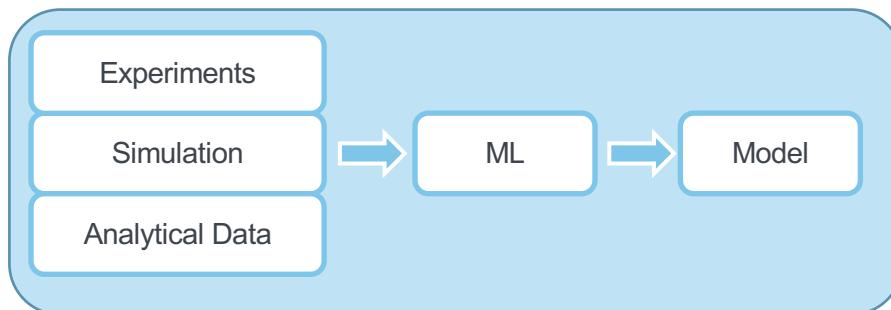
$$0 = R(U_h, x_h) + C(U_l, U_h) \approx \underbrace{R(U_h, x_h)}_{\text{"Closed" Equations at level h}} + \underbrace{M(U_h)}_{(2)}$$

# Machine Learning for data-driven submodels

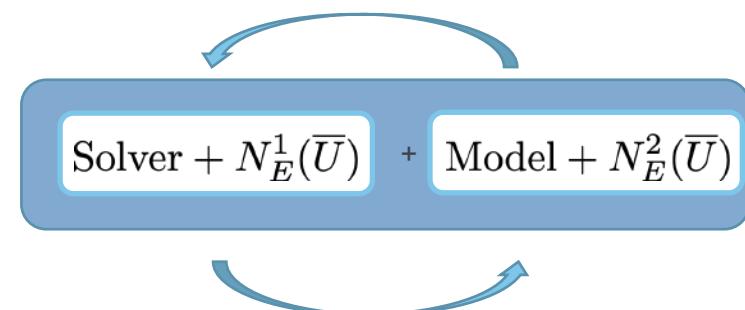
$$R(U_L, x_L, t_L) = -M(U_l, U_L)$$

$$R(U_L, x_L, t_L) \approx -\tilde{M}(U_L) \longrightarrow$$

- Riemann solver
- Forchheimer diffusion coefficient
- RANS closure coefficients
- ...



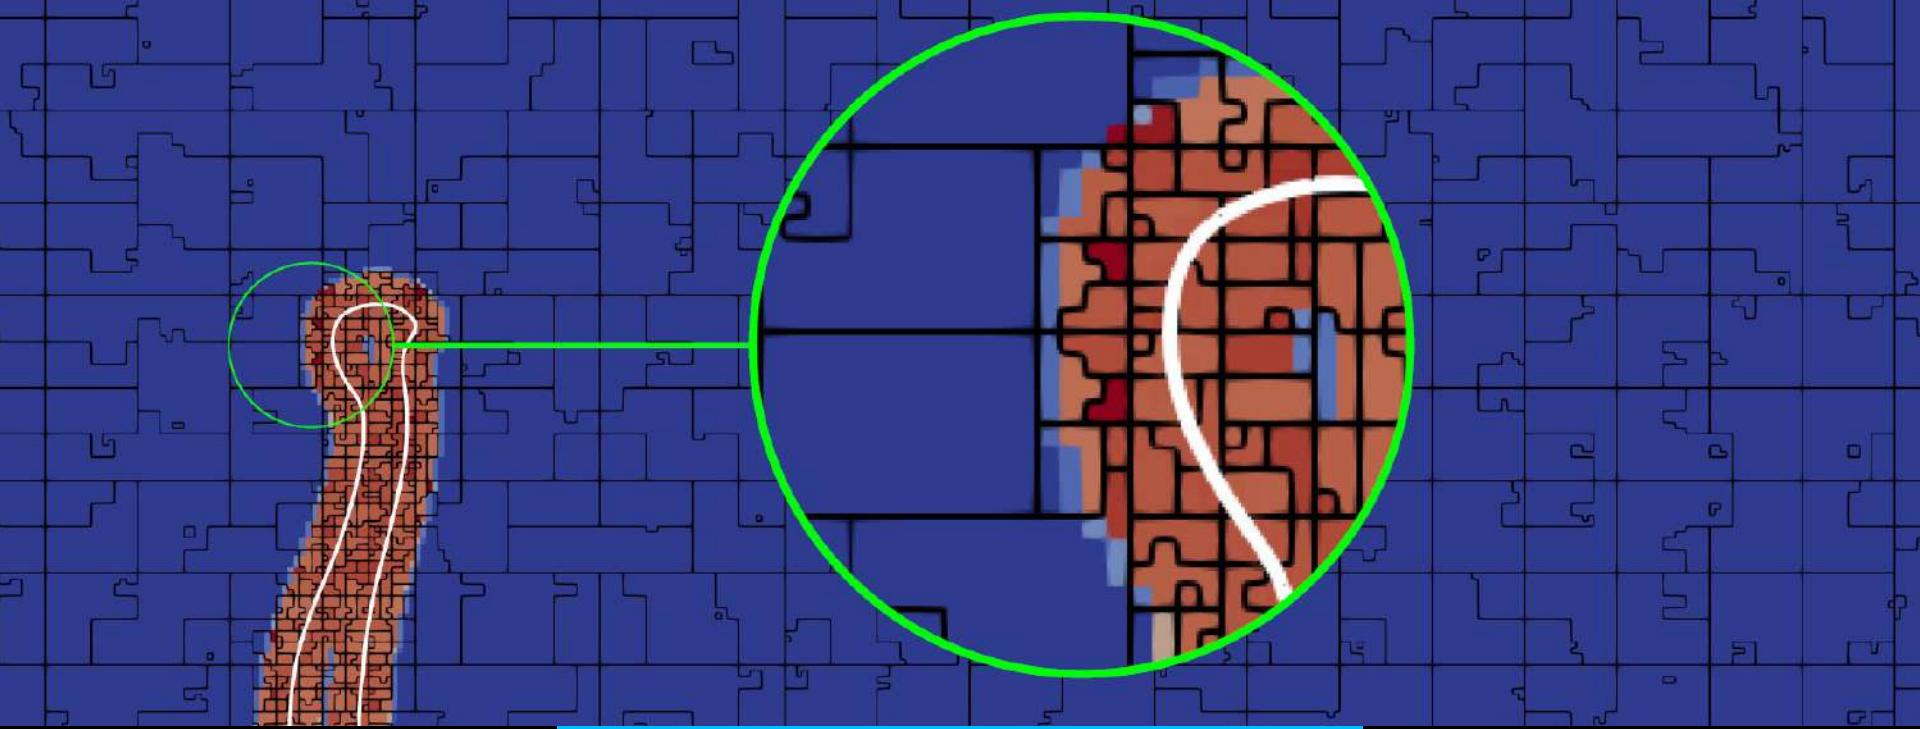
Stage 1: Offline training on well-defined data



Stage 2: Online inference on “real” data

# Small detour ahead: Our SL success stories

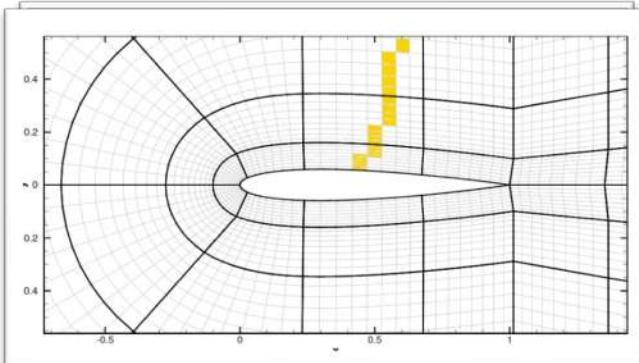
“ ”



# Data-Driven Local Shock Capturing for DG

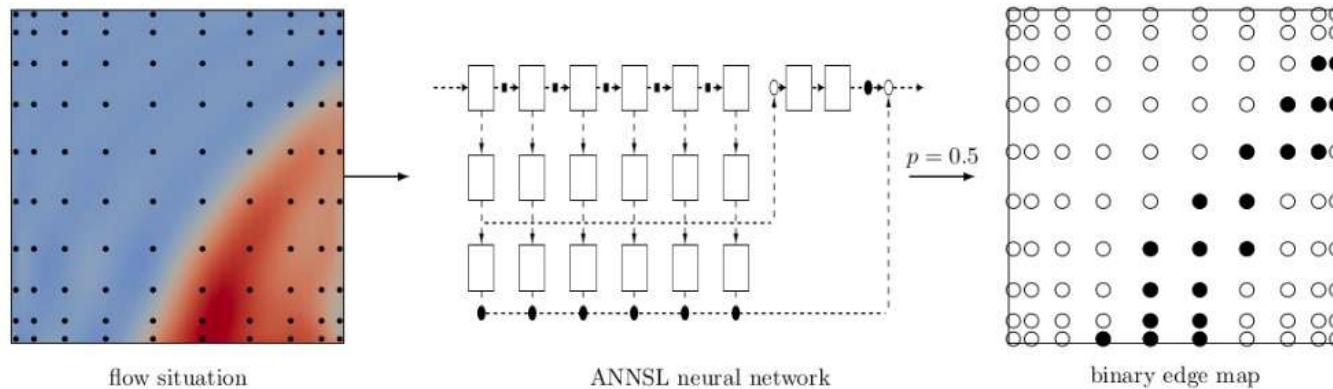
# Data-informed High Order Methods

- Discontinuities **conflict** with high order polynomials
- Stable numerical approximation through Shock Capturing: improves stability, decreases accuracy: use **sparingly!**
- First step: Detecting the occurrence of shocks: non-trivial, empiricism, **many parameters**
- For HO methods: Just detecting a “troubled cell” is not good enough: We need **localization on the element subscale**  $h/(N + 1)$



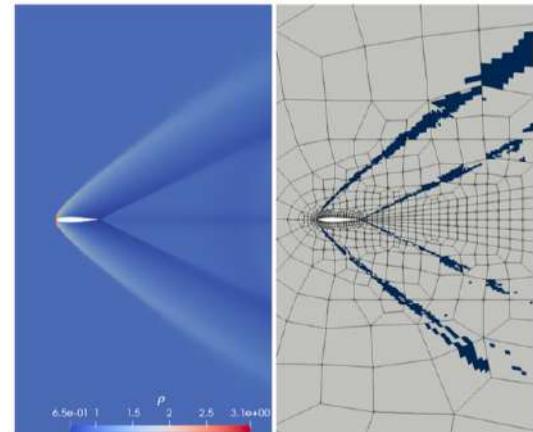
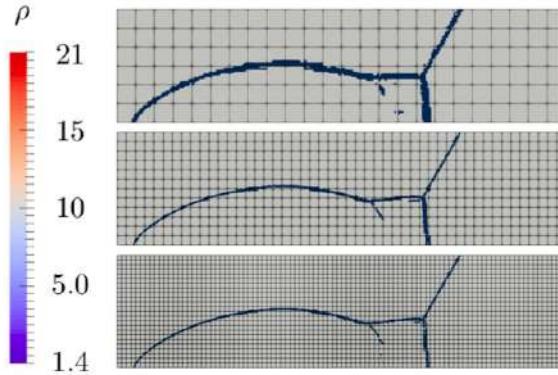
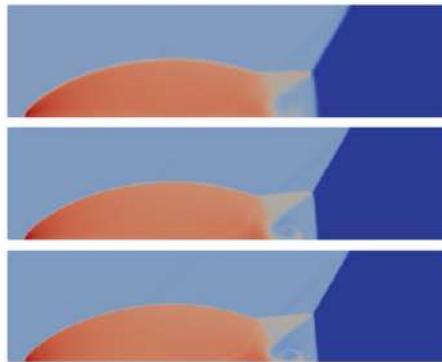
# Shock Locator

- Locate shock front **within high order element** on the nodal level
- Supervised learning strategy on **analytically labeled data**
- **Multiscale CNN networks**, adapted from edge detection



# Shock Locator

- Invariance w.r.t to grids, cases, etc: no retraining, no parameters
- Consistent, continuous shock front on varying resolutions: unstructured grids
- Localized, polynomial artificial viscosity with inner element shock capturing



# Localized AV for shock capturing

- Shu-Osher Shock Interaction

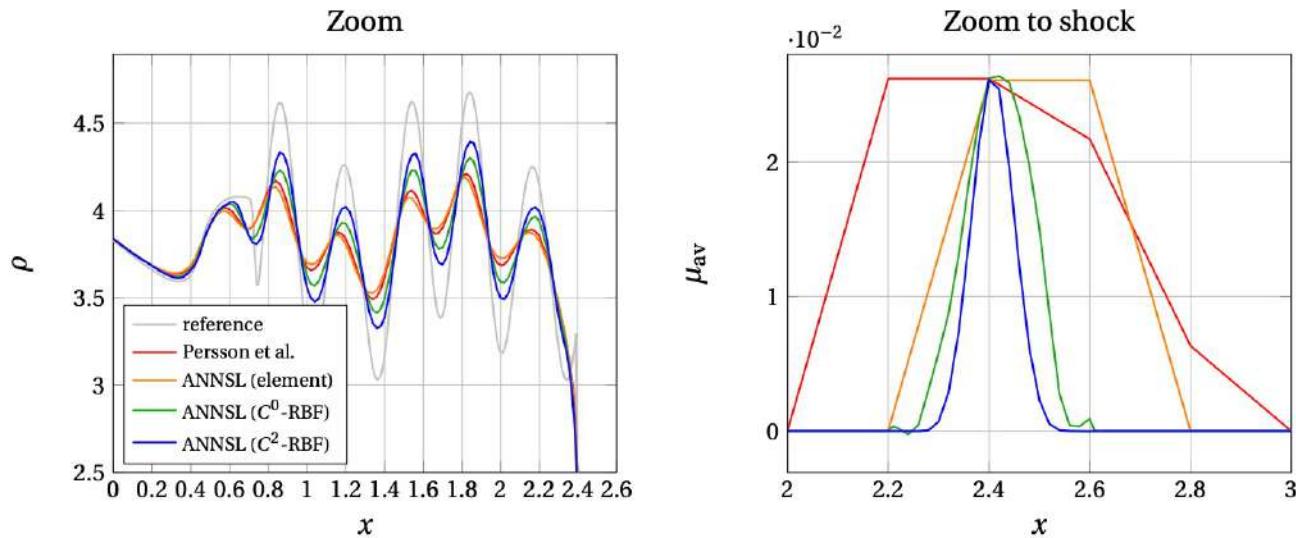


Figure 6: Density solution (left) and artificial viscosity (right) for Shu-Osher problem with  $N = 9$  and  $n_{\text{elems}} = 50$  at  $t = 1.8$ . The background grid in  $x$ -direction indicates the mesh spacing ( $\Delta x = 0.2$ ).

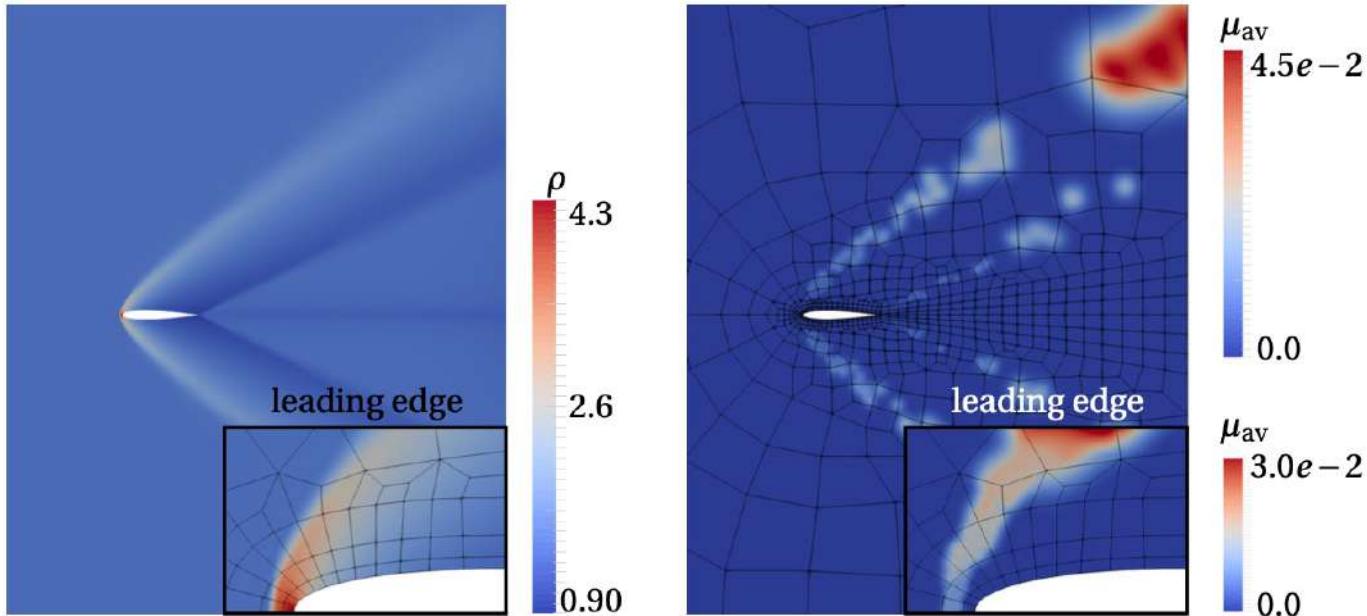


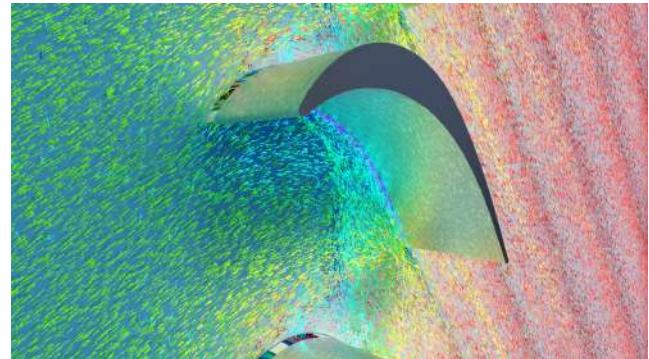
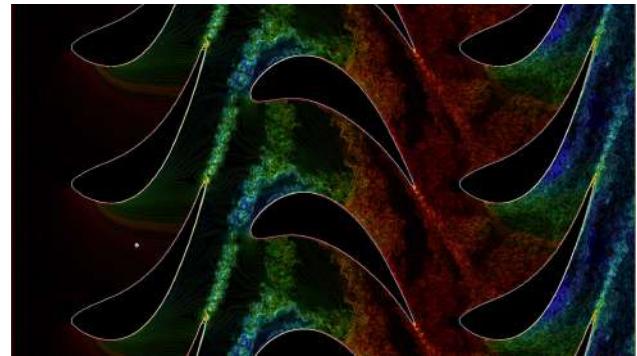
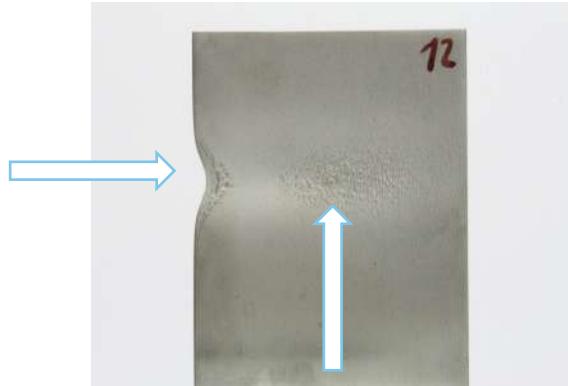
Figure 14: Density (left) and artificial viscosity with underlying computational mesh (right) of NACA0012 profile each with detailed view of leading edge at  $Ma = 2$  and  $t = 10$ , simulated with  $N = 5$ ,  $CFL = 0.5$ ,  $\sigma = 0.25$  and  $r_s = 0.75$ .



# Wall Models

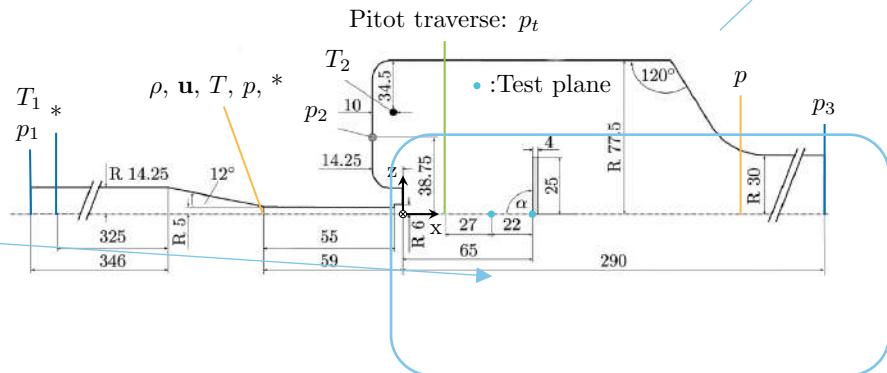
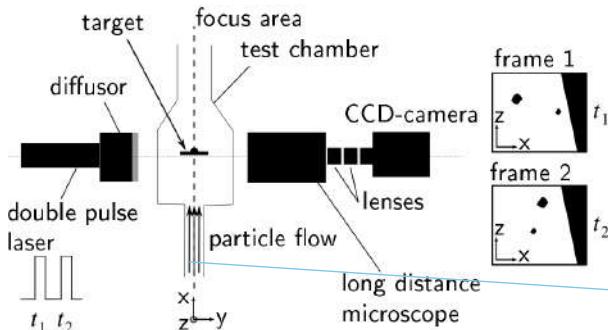
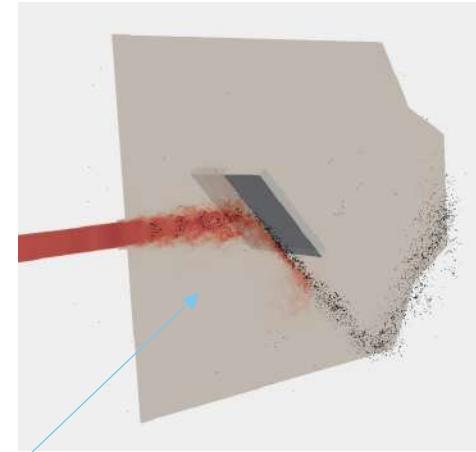
# Erosion in Turbomachines

- Erosion due to ingested particles limits lifespan and safety (Eyjafjallajökull, 2010)
- Turbulent flow – particle and **particle / wall interaction**
- Model **for individual P/W** needed: velocities, angles, breaking, rotation, etc.



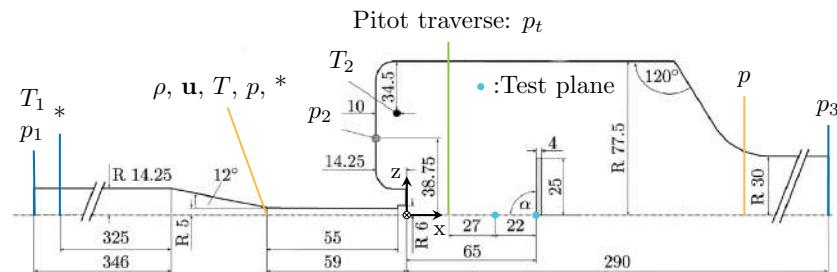
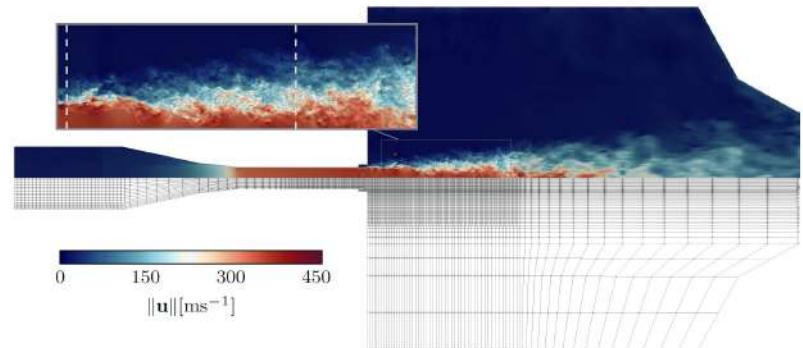
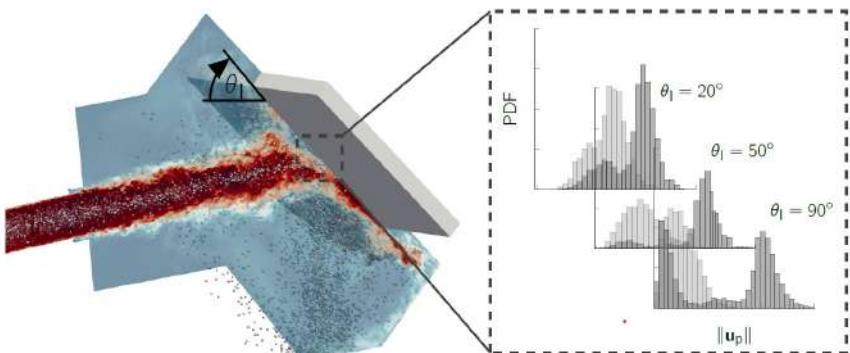
# Measurement Campaign

- Cooperation with Institute of Aircraft Propulsion Systems
- Problem: Particle paths are determined by **consecutive particle/wall interactions**
- Experimental and computational campaign
- Particles in turbomachines are **fast**

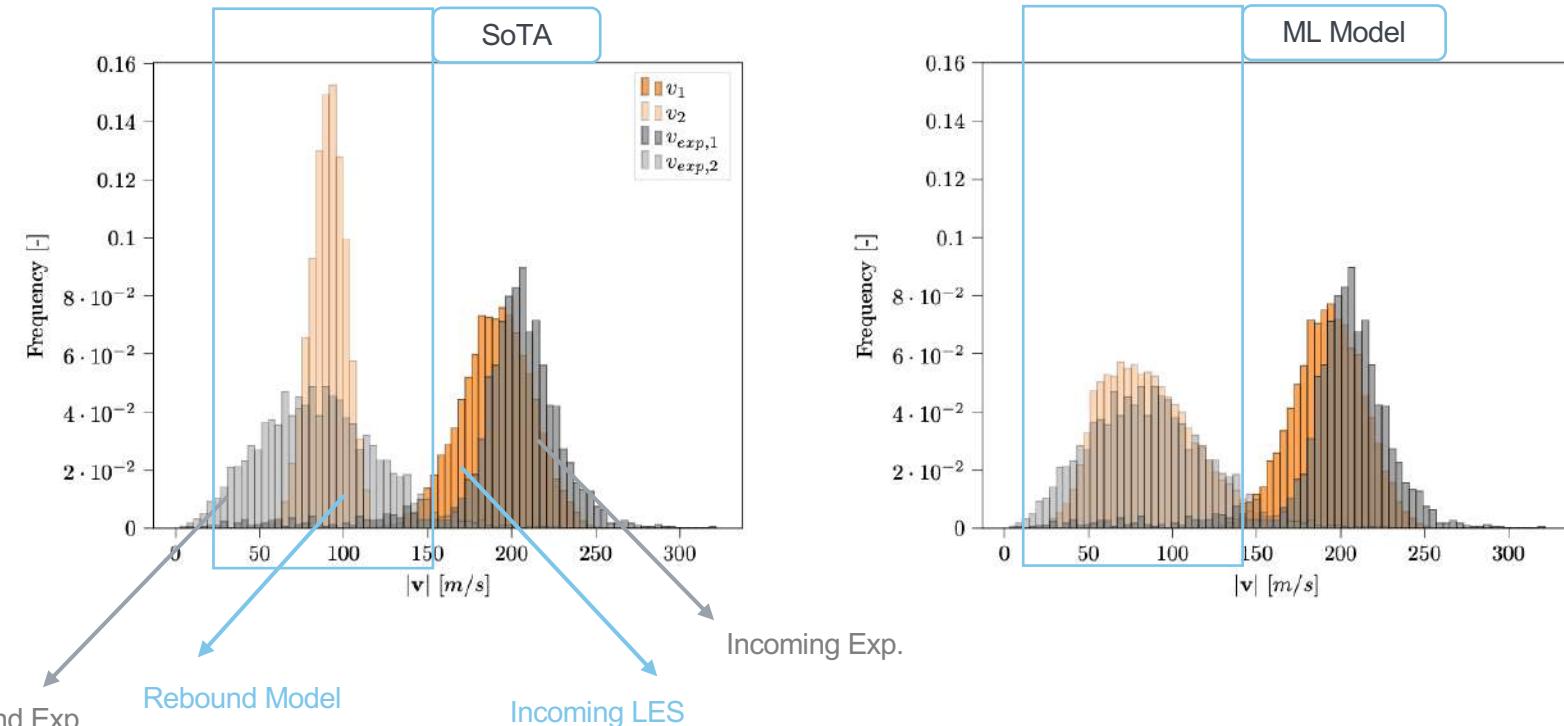


# Particle / Wall models from data

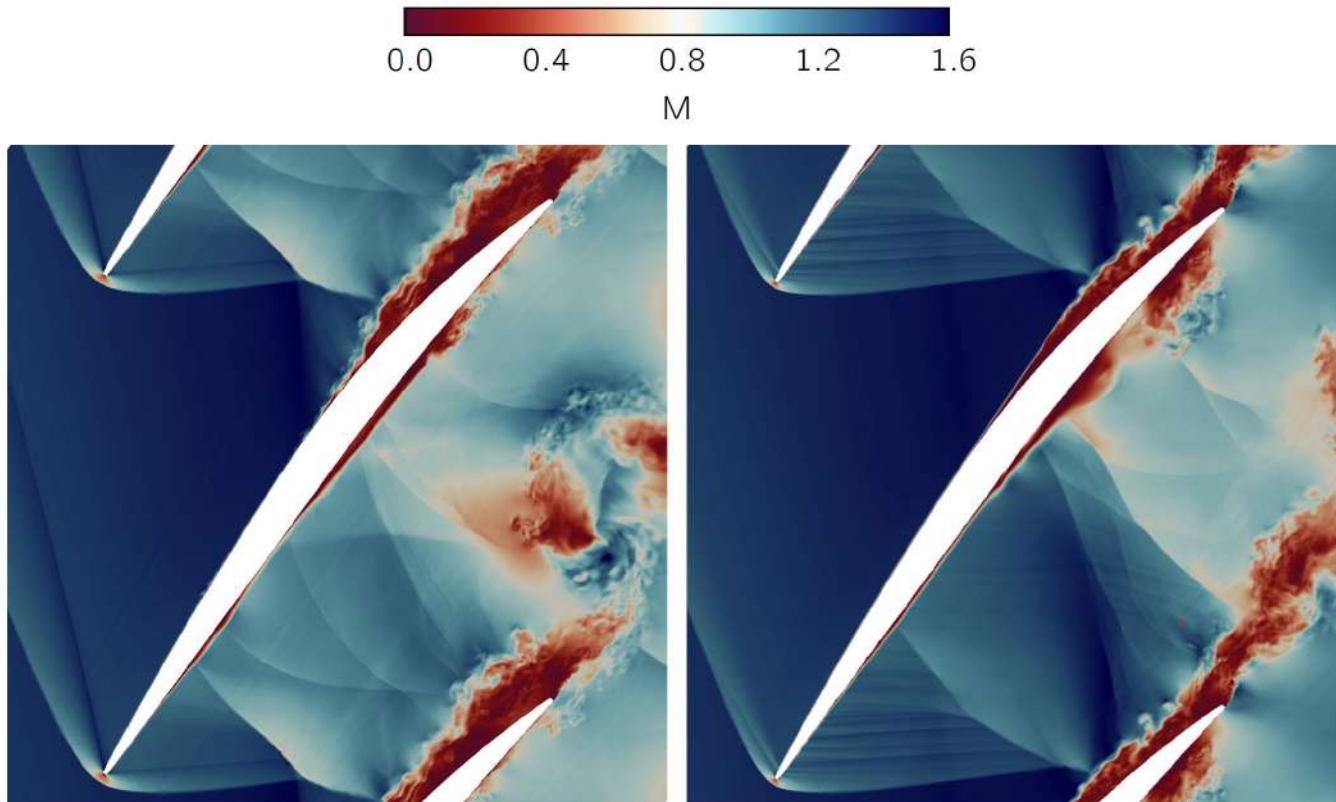
- Measurement in particle wind tunnel
- Supervised ML with a priori conditioning on physical constraints



# A posteriori model evaluation

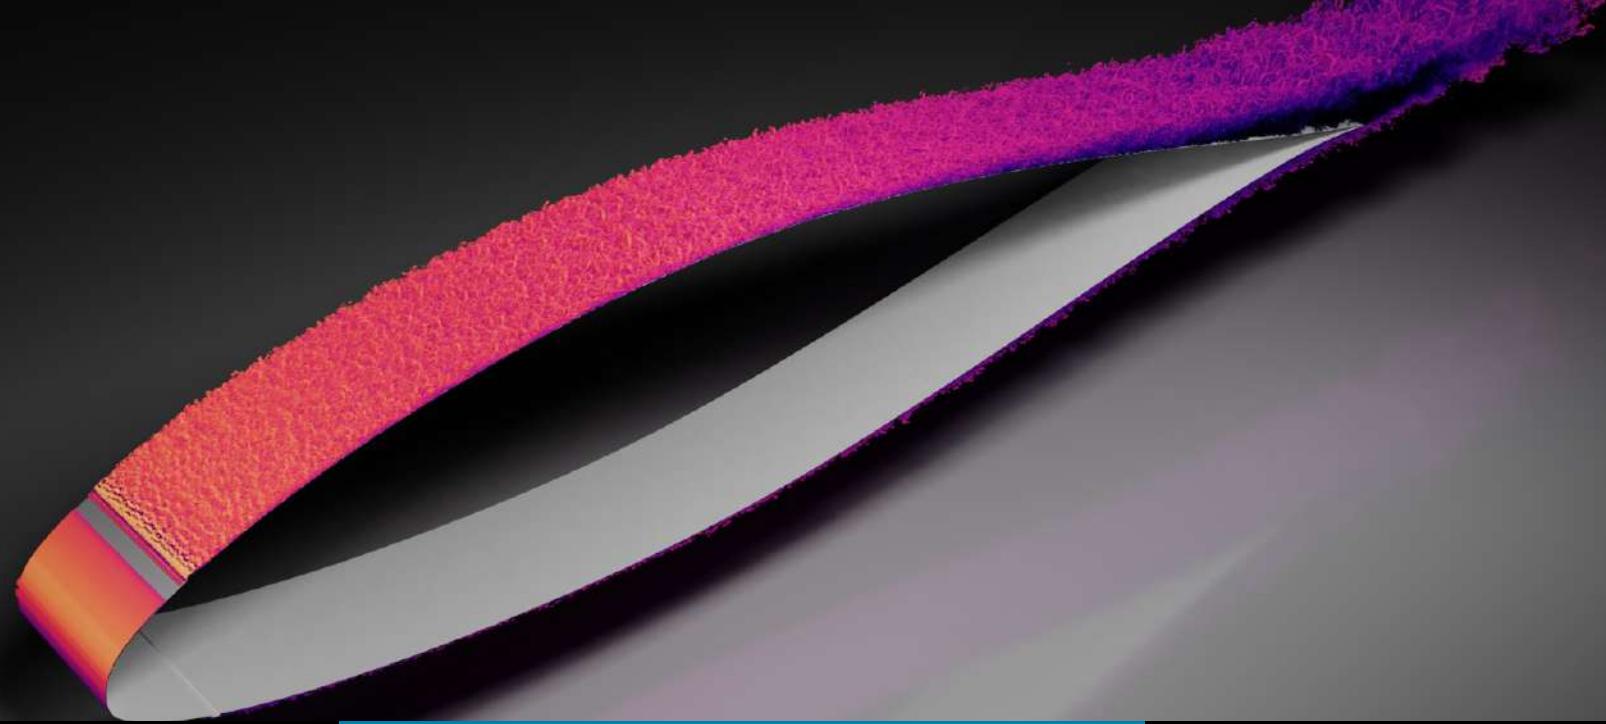


# ML-driven, time-resolved erosion simulations



# Back to LES

“ ”

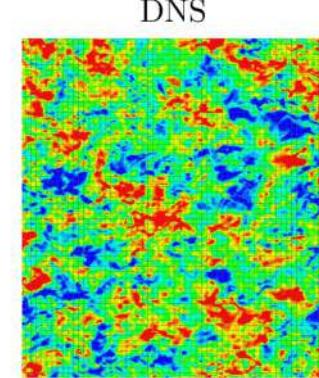


## Turbulence closures from data

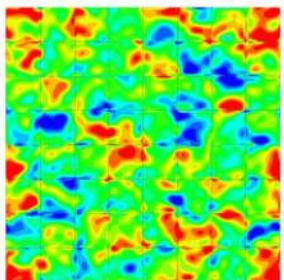
# The Scale Gap in Turbulence

$$\underbrace{R(U_l, x_l, t_l)}_{{\text{Governing Equations at level 1}}} = 0$$

$$\underbrace{R(U_L, x_L, t_L)}_{{\text{Governing Equations at level L}}} + \underbrace{M(U_l, U_L)}_{{\text{Influence of level 1 on L}}} = 0$$



Coarse Grid Data  $U$



Closure Terms  $\overline{R(F(U))}$

The Reynolds stresses:  
The footprint of the fine  
scale turbulence

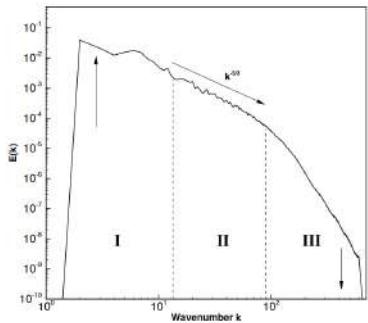
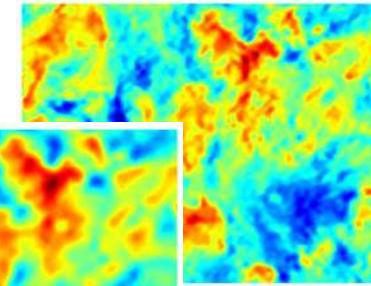
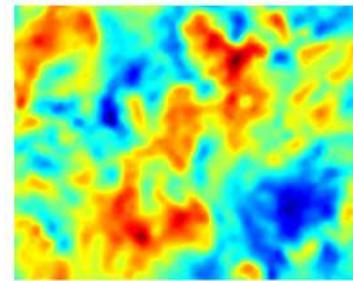
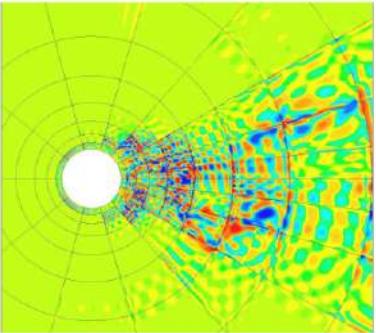
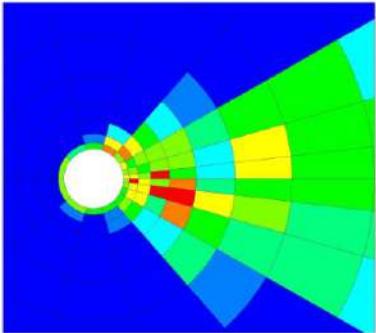
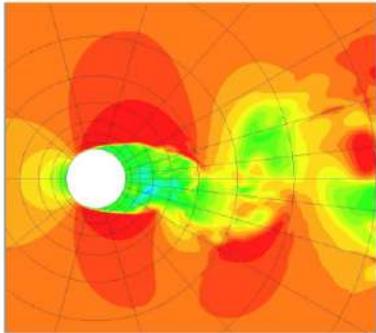
Data-driven models

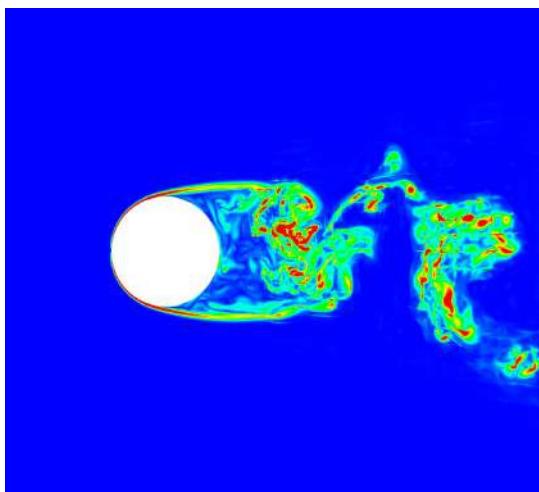
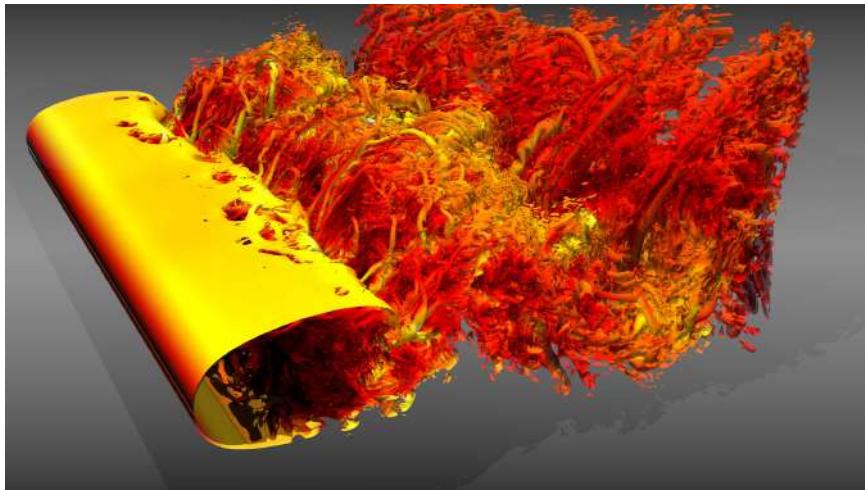
# LES - Basics

$$\overline{u(x)} = \int_{-\Delta x}^{\Delta x} u(r) G(x - r) dr$$

$$\frac{\partial \bar{u}_i}{\partial t} + \frac{\partial}{\partial x_j} (\bar{u}_i \bar{u}_j) = -\frac{1}{\rho} \frac{\partial \bar{p}}{\partial x_i} + 2\nu \frac{\partial}{\partial x_j} \bar{S}_{ij} - \frac{\partial \tau_{ij}}{\partial x_j}$$

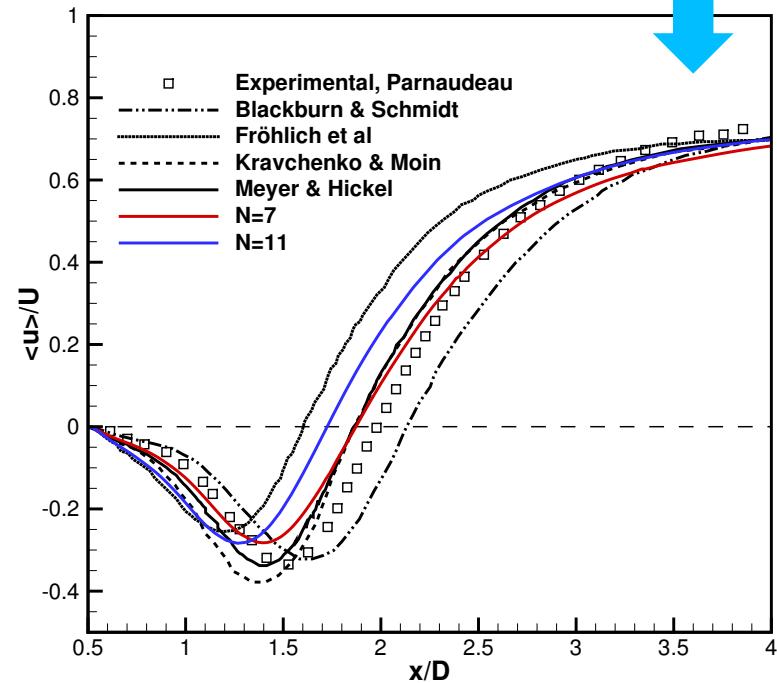
$$\overline{u_i u_j} = \tau_{ij} + \bar{u}_i \bar{u}_j$$



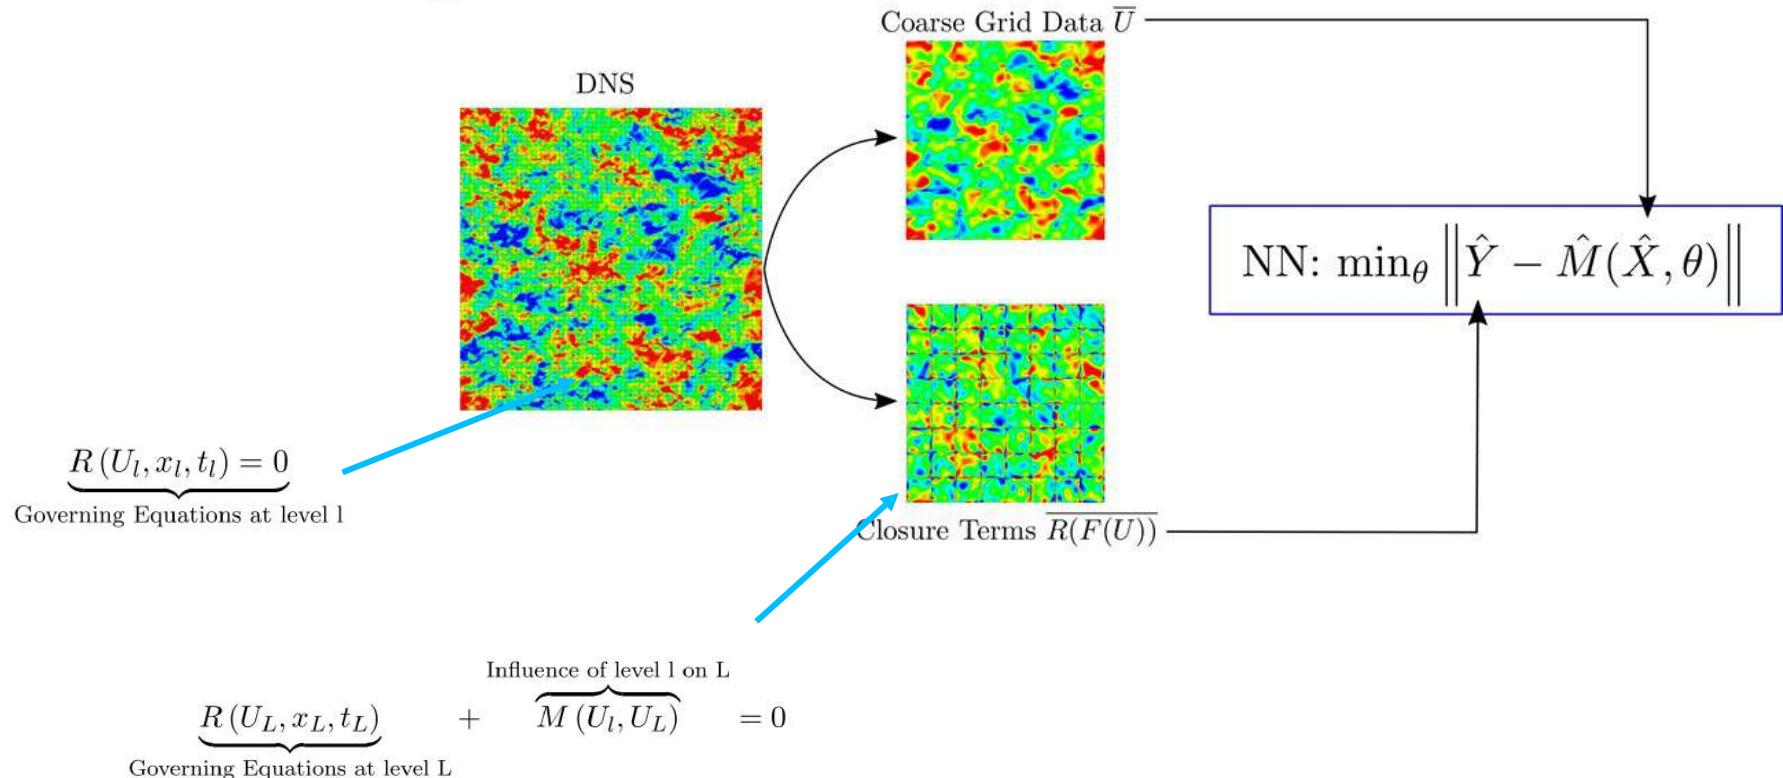


“It is the model’s fault!”

“Same equations, domains, models...”



# LES Closure Terms

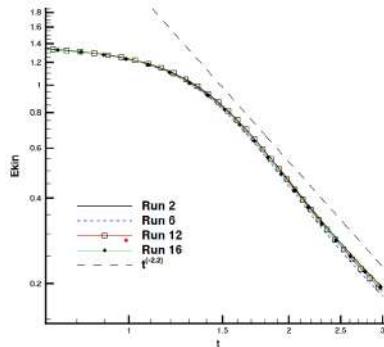
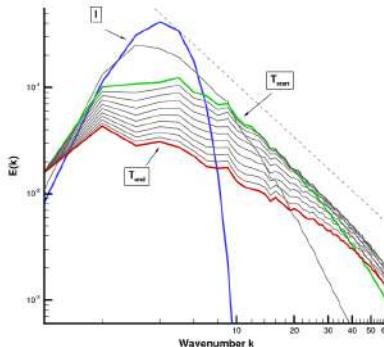


# Reynolds terms only!

“ ”

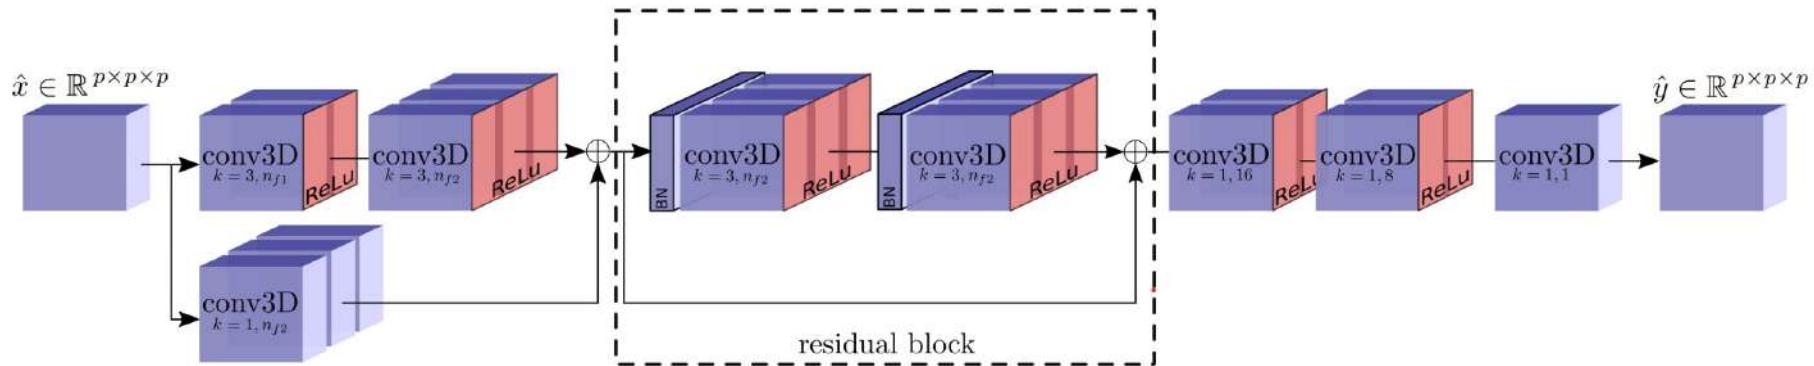
# Data Acquisition: Decaying Homogeneous Isotropic Turbulence

- Canonical test case of "turbulence in a box"
- **Ensemble of DNS runs** of decaying homogeneous isotropic turbulence with initial spectrum defined by Chasnov (1995) initialized by Rogallo (1981) procedure and  $Re\lambda = 180$  at start
- Data collection in the range of exponential energy decay: **25 DHIT** realizations with
- 134 Mio DOF each computed on CRAY XC40 (approx. 400,000 CPUh, 8200 cores)

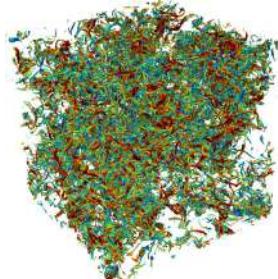
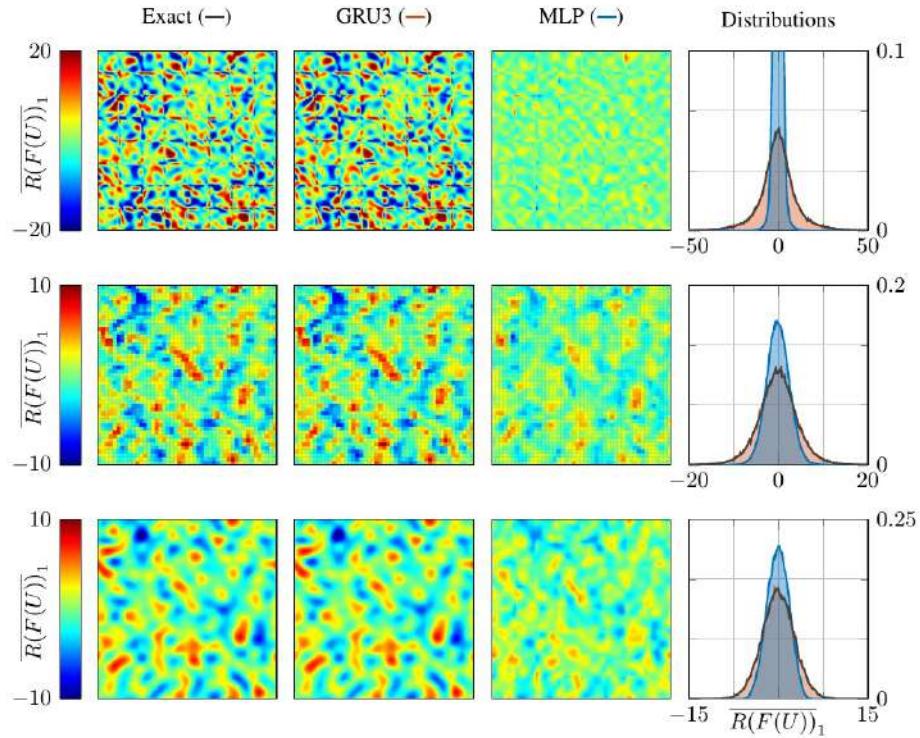


# Networks and Training

- CNNs with skip connections (RNN), batch normalization, ADAM optimizer, data augmentation
- Different network depths (no. of residual blocks)
- For comparison: MLP with 100 neurons in 1 hidden layer
- Implementation in Python / Tensorflow, Training on K40c and P100 at HLRS
- Split in training, semi-blind validation and blind test DHIT runs



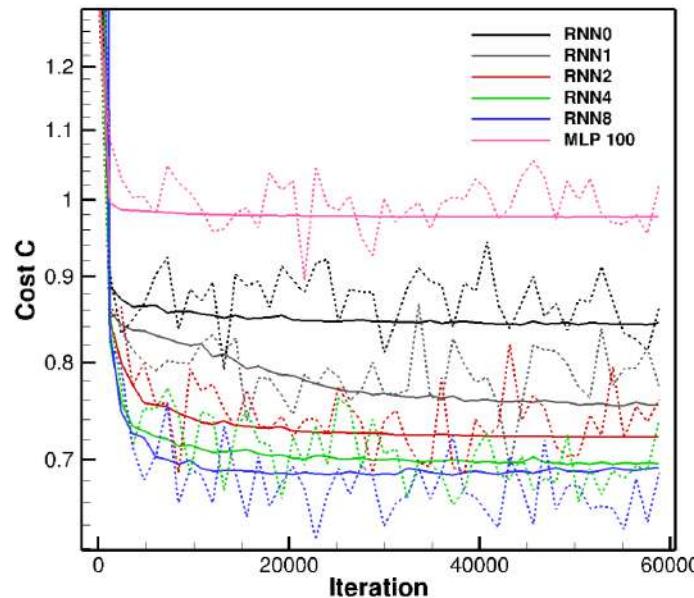
# LES Closure Terms



- 99.9999% Cross correlation with GRU models
- Near perfect accuracy in predicting Reynolds terms. Is this sufficient?

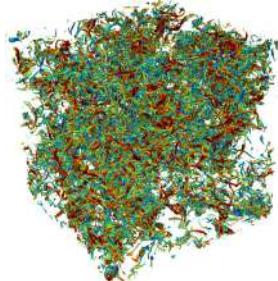
# Training Results I: Cost Function

- Cost function for different network depths
- RNNs outperform MLP, **deeper networks learn better**
- The approach is data-limited. NNs are very data-hungry.



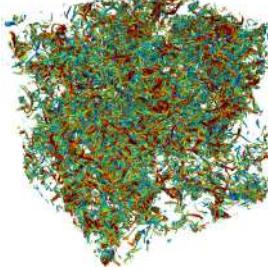
## Training Results II: Cross Correlation

Smagorinsky: about 0%,  
scale similarity models:  
about 60%

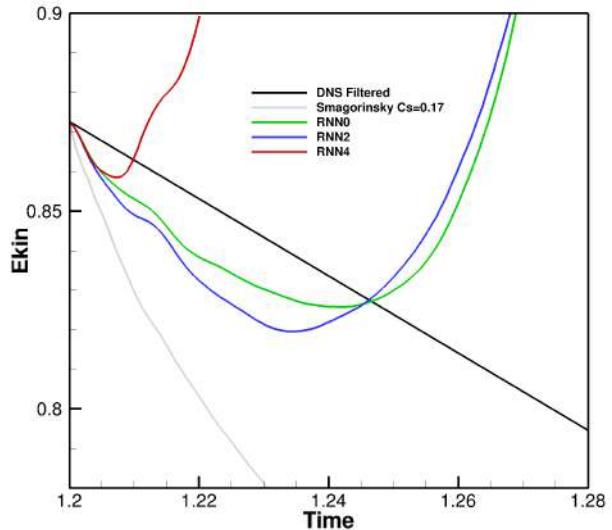
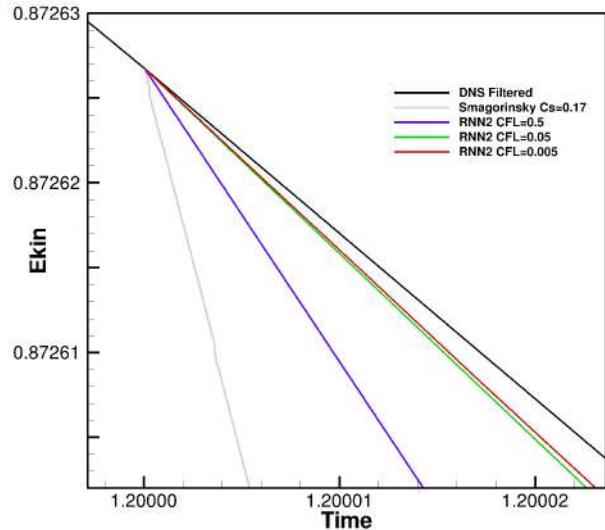


Network	$a, b$	$\mathcal{CC}(a, b)$	$\mathcal{CC}^{inner}(a, b)$	$\mathcal{CC}^{surf}(a, b)$
RNN0	$\overline{R(F(U))}^1, \overline{R(F(U))}^1^{ANN}$	0.347676	0.712184	0.149090
	$\overline{R(F(U))}^2, \overline{R(F(U))}^2^{ANN}$	0.319793	0.663664	0.134267
	$\overline{R(F(U))}^3, \overline{R(F(U))}^3^{ANN}$	0.326906	0.669931	0.101801
RNN4	$\overline{R(F(U))}^1, \overline{R(F(U))}^1^{ANN}$	0.470610	0.766688	0.253925
	$\overline{R(F(U))}^2, \overline{R(F(U))}^2^{ANN}$	0.450476	0.729371	0.337032
	$\overline{R(F(U))}^3, \overline{R(F(U))}^3^{ANN}$	0.449879	0.730491	0.269407

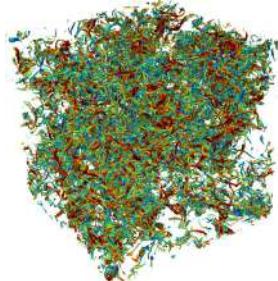
# LES with trained ANN model



- Direct closure, **no constraints incorporated**
- No long term stability, but short term stability and dissipation
- Improvement: Loss function penalization of anti-dissipative stress tensor



# LES Closure Terms: GRU



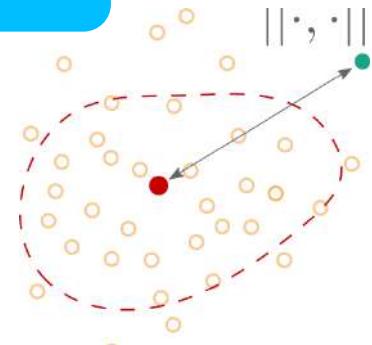
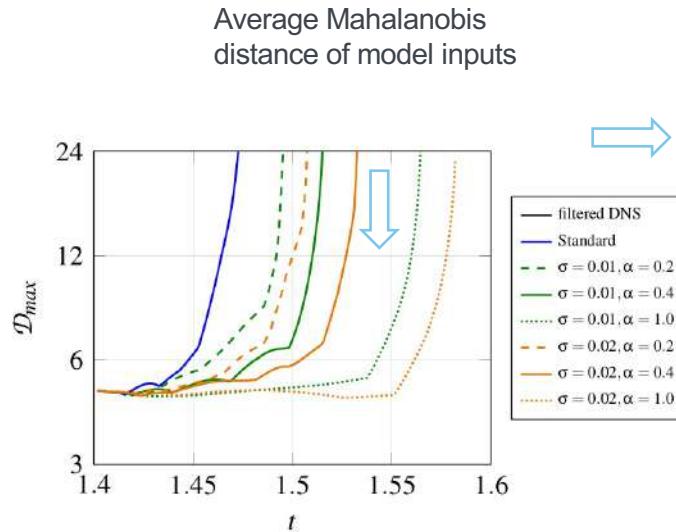
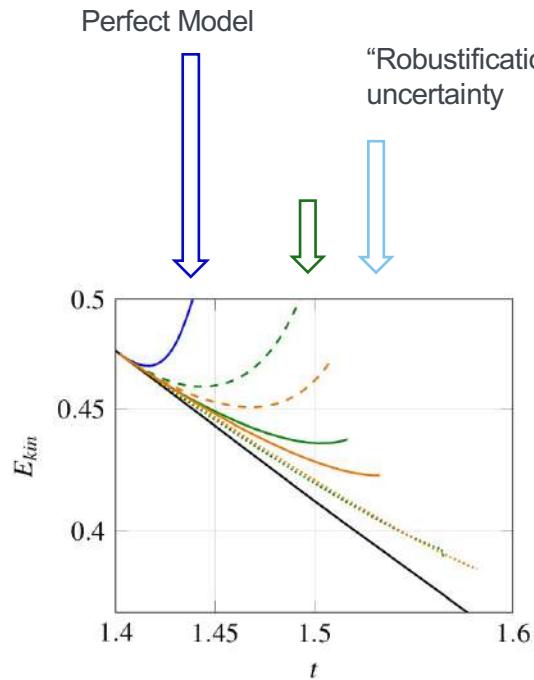
Network	# Parameter	Time (GPU)	Time (CPU)	$L_2$ -Error	CC
MLP	6,720	6 ms	28 ms	$3.0 \cdot 10^{+1}$	66.0%
CNN	187,088	72 ms	198 ms	$2.1 \cdot 10^{+1}$	78.7%
LSTM (3 $\Delta t$ )	39,744	62 ms	340 ms	$1.3 \cdot 10^{-1}$	99.9%
GRU (3 $\Delta t$ )	31,578	59 ms	319 ms	$1.1 \cdot 10^{-1}$	99.9%

$$\frac{\partial \bar{U}}{\partial t} + \widetilde{R}((\bar{U})) = \widetilde{R}((\bar{U})) - \overline{R((U))}$$

Prediction on test set with arbitrary accuracy > 99.9999% CC

# LES Closure Terms: Direct Approach

Chaotic, dynamical systems and MDI are not a good match!



- Arbitrarily accurate closure terms, but stability is lacking: **Stability training and incorporating of physical constraints helps only a bit**
- Can we train robust and accurate models this way for practical applications?

# LES Closure Terms: Useful Models

Predict Reynolds stresses with ML, then find best fit to stable model

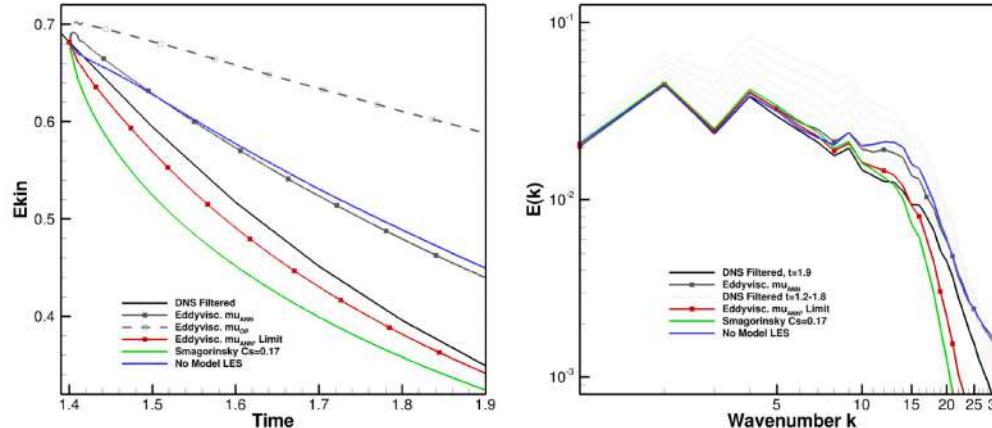
$$\frac{\partial \overline{U}}{\partial t} + \tilde{R}(F(\overline{U})) = \underbrace{\tilde{R}(F(\overline{U})) - \overline{R(F(U))}}_{\text{data-based eddy viscosity model}}.$$

Accurate Prediction, stability issues

- Simplest model: **Eddy viscosity approach** with  $\mu_{ANN}$  from

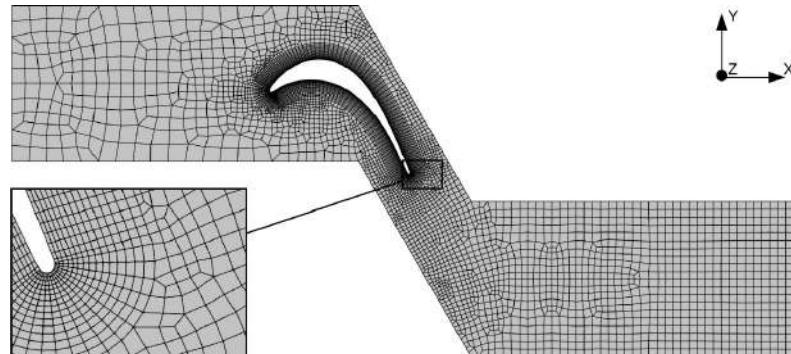
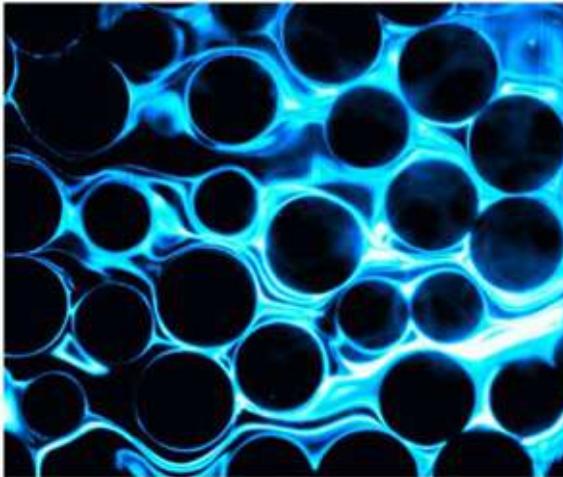
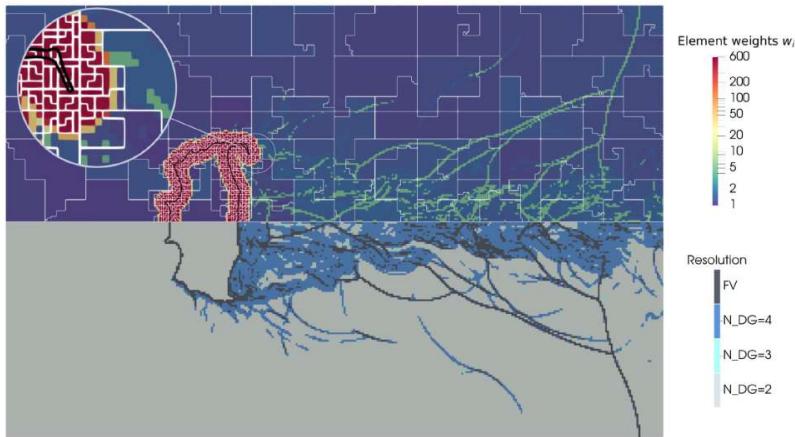
$$\tilde{R}(F(\overline{U^i})) - \overline{R(F(U^i))} \approx \mu_{ANN} \tilde{R}(F^{visc}(\overline{U^i}, \nabla \overline{U^i}))$$

Project prediction onto guaranteed stable model



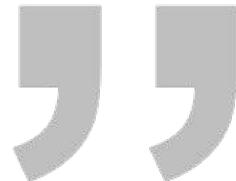
# We want LES models that work for ...

- Adaptive discretizations
- Unstructured grids
- All kinds of flows across  $Re$ ,  $Ma$ ,  $Sc$ , ...

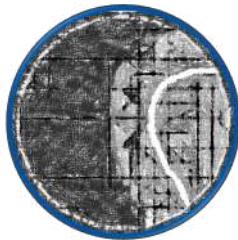


# Maybe SL is not the way forward?

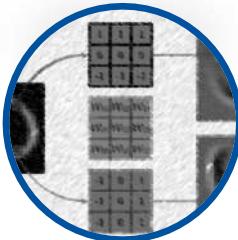
1. Do “smarter” SL
2. Do “smarter” ML



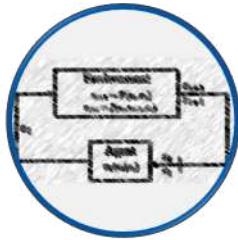
# Overview



Part 1: Multi-X problems and PDE solvers

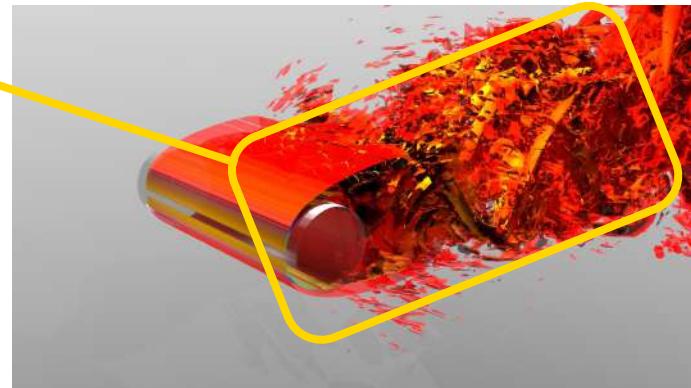
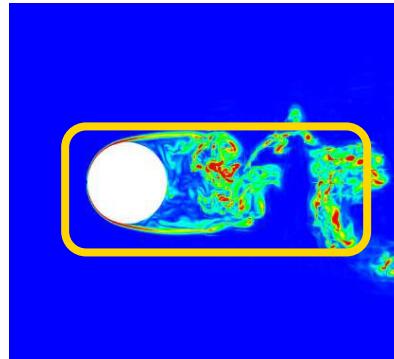
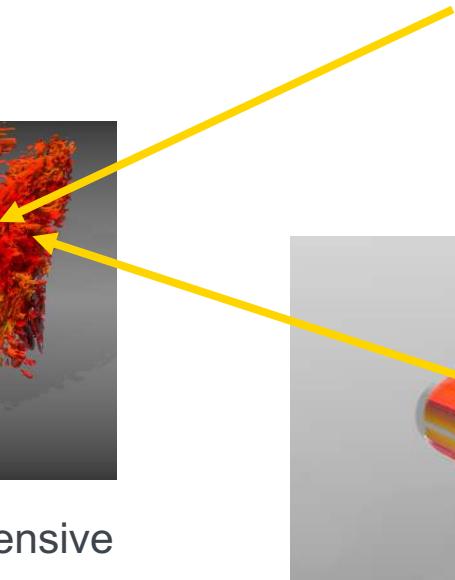
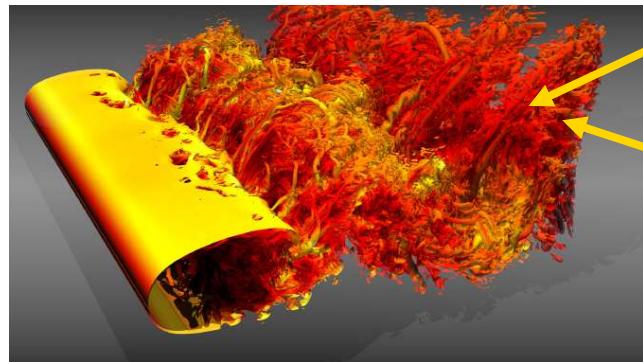


Part 2: Supervised learning for PDE submodels



Part 3: From data-driven to integrated CFD/ML

# Revisiting the LES problem



- Computationally too expensive
- **Data rich and efficiency poor?**
- SL Database and Curation?
- Closure models must account for **physical** and **computational** subfilter stresses

# Large Eddy Simulation - Definition

$$u = \bar{u}(x, t) + u(x, t)' \Rightarrow \bar{u}(x, t) = \int_{\Delta r} u(r, t) G(x - r) dr$$

- Separate  $\Delta r$  and  $h$
- Explicit filtering:  $h \rightarrow 0$
- Discretization scheme not relevant
- **Caveat: homogeneity and isotropy, boundary conditions, realizability, commutation...**

0.5% on google scholar

- Joined  $\Delta r$  and  $h$
- Implicit filtering:  $h \not\rightarrow 0$
- Discretization scheme defines filter kernel
- $\Leftarrow$  **plus discretization parameters and errors**

99.5% on google scholar, 100% for “industrial” LES

We derive the LES equations and the closure terms for a convolution with a known, homogenous, linear filter kernel with commutation

$$\bar{\psi}(x) := F(\psi(x)) = \int_{-\infty}^{\infty} k(\xi - x, \Delta) \psi(\xi) d\xi,$$

Use Burgers equations as the simplest model for non-linearity

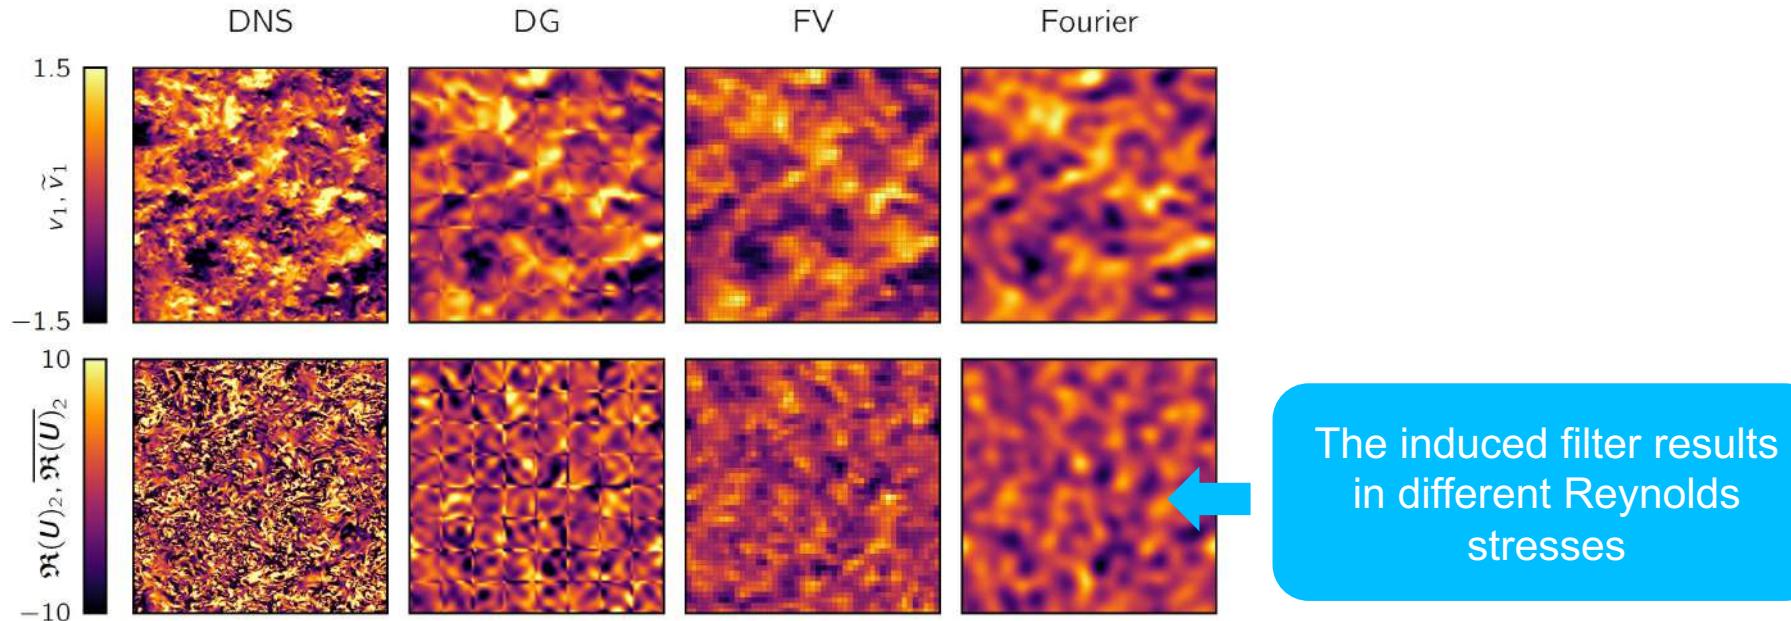
$$R(u) = \frac{\partial u}{\partial t} + \frac{1}{2} \frac{\partial u^2}{\partial x} = 0,$$

$$\begin{aligned} \overline{R(u)} &= \frac{\partial \bar{u}}{\partial t} + \frac{1}{2} \frac{\partial \bar{u}^2}{\partial x} = \frac{1}{2} \underbrace{\left( \frac{\partial \bar{u}^2}{\partial x} - \frac{\partial \bar{u}^2}{\partial x} \right)}_{C_1[F; \frac{\partial}{\partial x}](u^2)} \\ &= \frac{\partial \bar{u}}{\partial t} + \frac{1}{2} \frac{\partial \bar{u} \bar{u}}{\partial x} = C_1 + \frac{1}{2} \underbrace{\left( \frac{\partial \bar{u} \bar{u}}{\partial x} - \frac{\partial \bar{u}^2}{\partial x} \right)}_{C_2[F; ()^2](u)}, \end{aligned}$$

- You find this as the Reynolds stresses in any textbook

Almost all practical LES: filtering is implicit, through grid / operator / discretization:

1.) Discretization characteristics **induce filter shape**



2.) Resolved and **represented** scales

$$h \approx \Delta$$

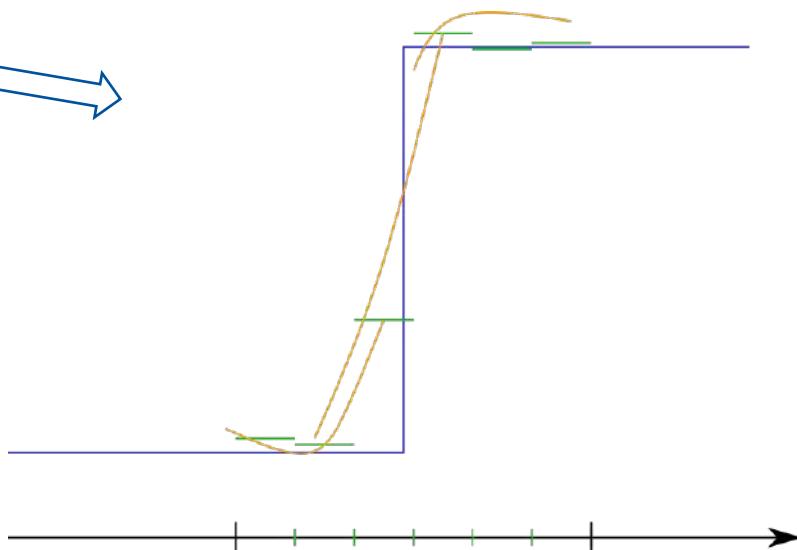
Discrete flux operator induces **computational subgrid stresses**

$$\overline{R(u)} = \frac{\partial \bar{u}}{\partial t} + \frac{1}{2} \frac{\delta \bar{u} \bar{u}}{\delta x} = C_1 + C_2 + \underbrace{\frac{1}{2} \left( \frac{\delta \bar{u} \bar{u}}{\delta x} - \frac{\partial \bar{u} \bar{u}}{\partial x} \right)}_{C_3[\partial_x; \delta_x]},$$

$$\overline{R(u)} = \frac{\partial \bar{u}}{\partial t} + \frac{1}{2} \frac{\delta \bar{u} \bar{u}}{\delta x} = C_1 + C_2 + \frac{1}{2} \frac{\partial}{\partial x} \underbrace{(\widehat{\bar{u} \bar{u}} - \bar{u} \bar{u})}_{C_3^*[\widehat{\cdot}; \cdot]}; \delta_x f = \partial_x \hat{f}$$

$$\overline{R(u)} = \frac{\partial \bar{u}}{\partial t} + \frac{1}{2} \frac{\delta \bar{u} \bar{u}}{\delta x} = C_1 + C_2 + \frac{1}{2} \left( \underbrace{\frac{\delta \bar{u} \bar{u}}{\delta x}}_{C_3^*[:,]} - \underbrace{\frac{\partial \bar{u} \bar{u}}{\partial x}}_{\cdot} \right),$$

$$\overline{R(u)} = \frac{\partial \bar{u}}{\partial t} + \frac{1}{2} \frac{\delta \bar{u} \bar{u}}{\delta x} = C_1 + C_2 + \frac{1}{2} \frac{\partial}{\partial x} \underbrace{(\widehat{u} \bar{u} - \bar{u} \bar{u})}_{C_3^*[:,]}; \delta_x f = \partial_x \hat{f}$$



Linear numerical schemes for solving partial differential equations, having the property of not generating new extrema (monotone schemes), can be at most first-order accurate.

Think: Upwinding schemes, characteristic based schemes, WENO, SUPG, ...

# Challenges for data-driven modelling

$$\overline{R(u)} = \frac{\partial \overline{u}}{\partial t} + \frac{1}{2} \frac{\delta \overline{u} \overline{u}}{\delta x} = C_1 + C_2 + \frac{1}{2} \underbrace{\left( \frac{\delta \overline{u} \overline{u}}{\delta x} - \frac{\partial \overline{u} \overline{u}}{\partial x} \right)}_{C_3[\partial_x; \delta_x]},$$

Make SL models scale-aware

Train SL models on Reynolds stresses from DNS or Exp.

$$\delta^n = \delta(\overline{u}^n), \quad \overline{u}^{n+1} = \overline{u}(\delta^n)$$

Knowable a priori

Not knowable a priori

“... the computational large-eddy closure typically contains an important contribution that is sensitive to the adopted spatial discretization”

“... considerable improvements of practical large-eddy simulation may be achieved in case the closure problem in the modified equation is made to explicitly account for the coarseness as well as the type of discretization.”

PHYSICS OF FLUIDS 17, 125103 (2005)

#### Numerically induced high-pass dynamics in large-eddy simulation

Bernard J. Geurts<sup>a)</sup>

*Multiscale Modeling and Simulation, Numerical Analysis and Computational Mechanics,  
Department of Applied Mathematics, University of Twente, P.O. Box 217, 7500 AE Enschede,  
The Netherlands and Anisotropic Turbulence, Faculty of Applied Physics,  
Eindhoven University of Technology, P.O. Box 513, 5600 MB Eindhoven, The Netherlands*

Federik van der Bos

*Numerical Analysis and Computational Mechanics, Department of Applied Mathematics, Faculty EEMCS,  
University of Twente, P.O. Box 217, 7500 AE Enschede, The Netherlands*

We derive the LES equations and the closure terms for a known, homogenous, linear filter kernel with commutation...

... Scale Separation filtering in practical LES is done implicitly with a partially unknown, inhomogeneous, possibly anisotropic and often non-linear filter kernel

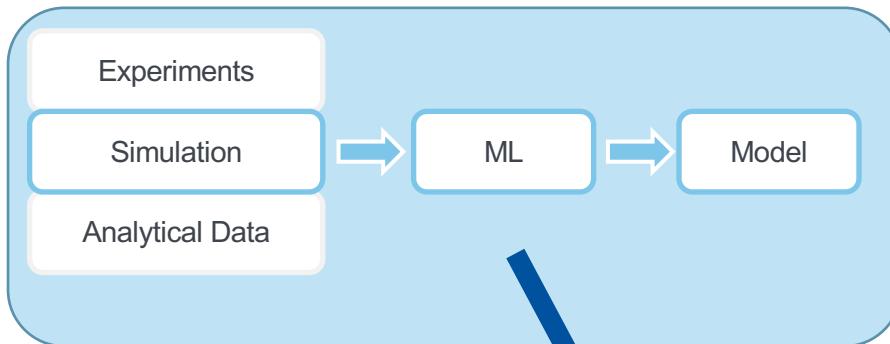
**What makes LES so  
tricky? more  
nonlinearities than we  
thought!**

”

## Challenge: A simple example

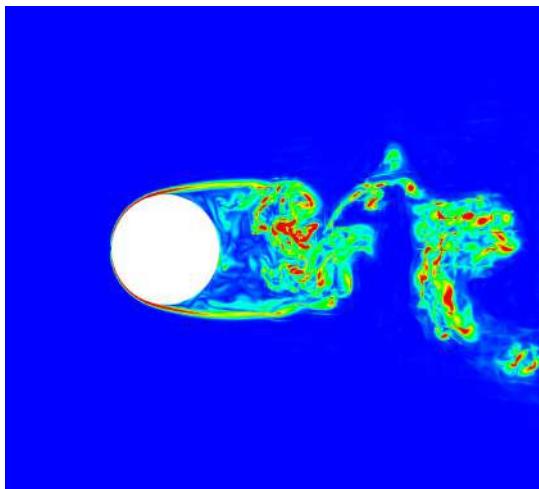
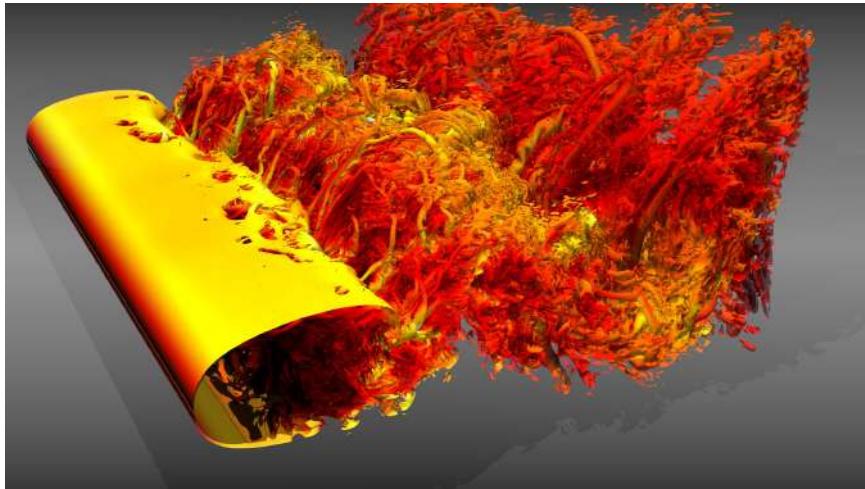
$$\frac{\partial u}{\partial t} + \frac{1}{2} \frac{\partial u^2}{\partial x} = \nu \frac{\partial^2 u}{\partial x^2}$$
$$\frac{\partial u}{\partial t} + R(U) = L(U)$$

$$\frac{\partial \bar{u}}{\partial t} + R(\bar{U}) = L(\bar{u}) + \underbrace{R(\bar{u}) - \bar{R(u)}}_S$$



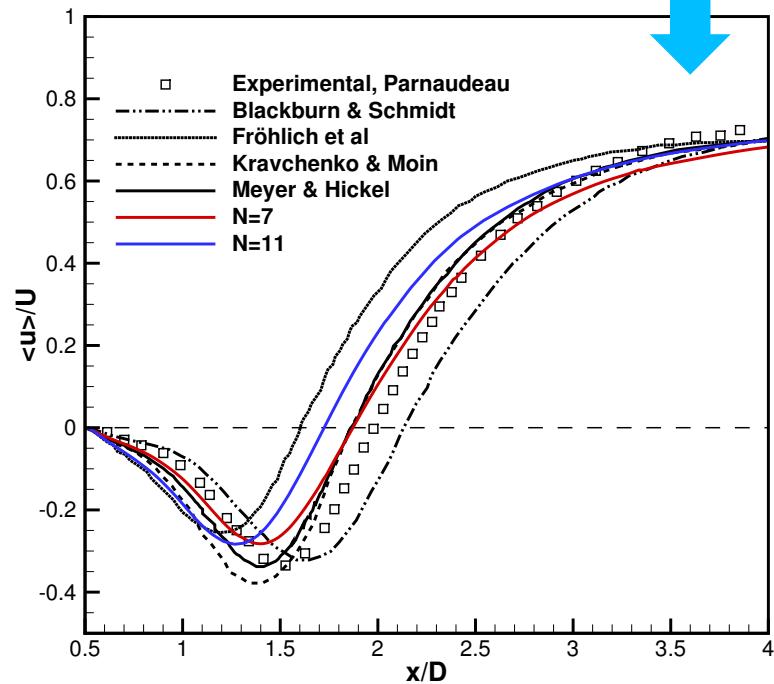
$$S \approx \frac{\partial}{\partial x} \left( \nu_{eddy} \frac{\partial \bar{u}}{\partial x} \right)$$

- Task: Find an optimal eddy viscosity for the Burgers equation from DNS data through SL and then use is **in your favorite numerical scheme**
- Beware: The gradient operator for an FD scheme, FV scheme, FE scheme looks different on the coarse grid!

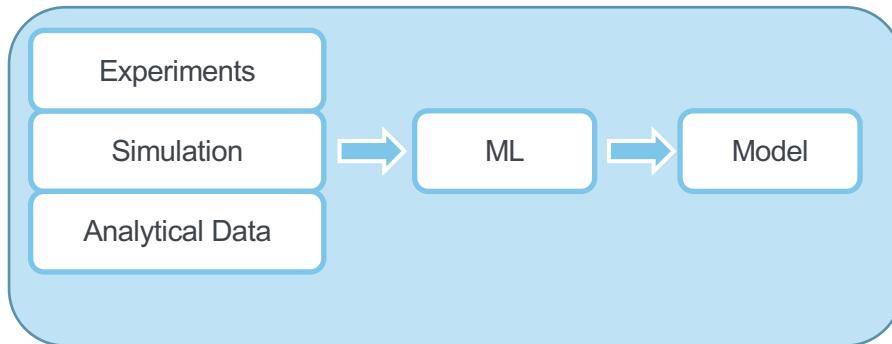


“Is it the model’s fault? Did we account for C1 and C3?”

“Same equations, domains, models...”

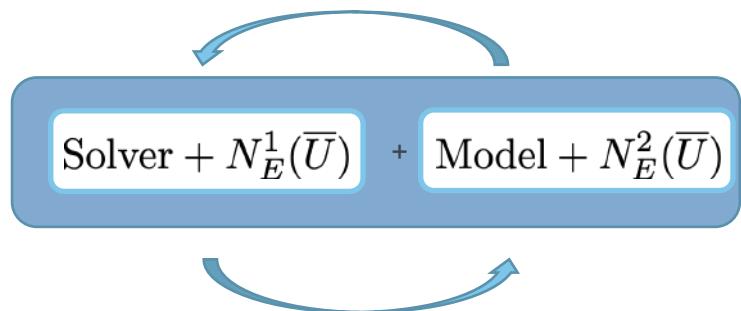


# Learning for dynamical systems: SL



Stage 1: Offline training on well-defined data

- Careful data generation and pre-processing
- Data analysis as a first step
- Dimensionality reduction
- Outlier detection
  - ML model formulations with guarantees (symmetries, statistics, equivariances, scaling, ...)
  - KEP networks
  - Cost function penalization
  - Robustness against uncertain inputs



Stage 2: Online inference on “real” data

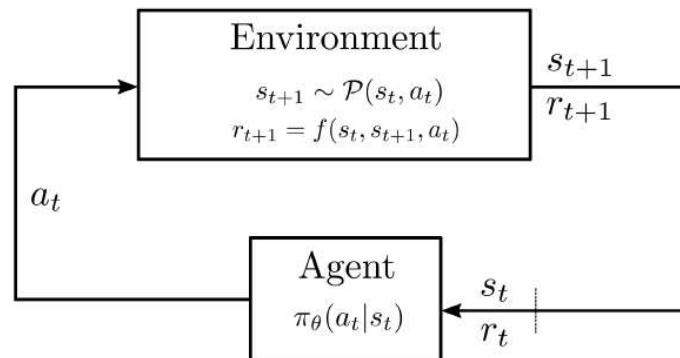
- Self-aware models with safe fallbacks
- Transfer learning
- Intermittent re-training

# Learning for dynamical systems: RL

- Simulations deal with (discrete) dynamical systems
- Submodels need to be **consistently integrated** with the discretization: a simple gradient looks very different for a FD, FV, FE scheme
- Supervised learning has two problems: How do we define **a lot of true data and make models robust?**
- RL is a different paradigm from ML: **Learning to take optimal actions in a dynamical system from experience**

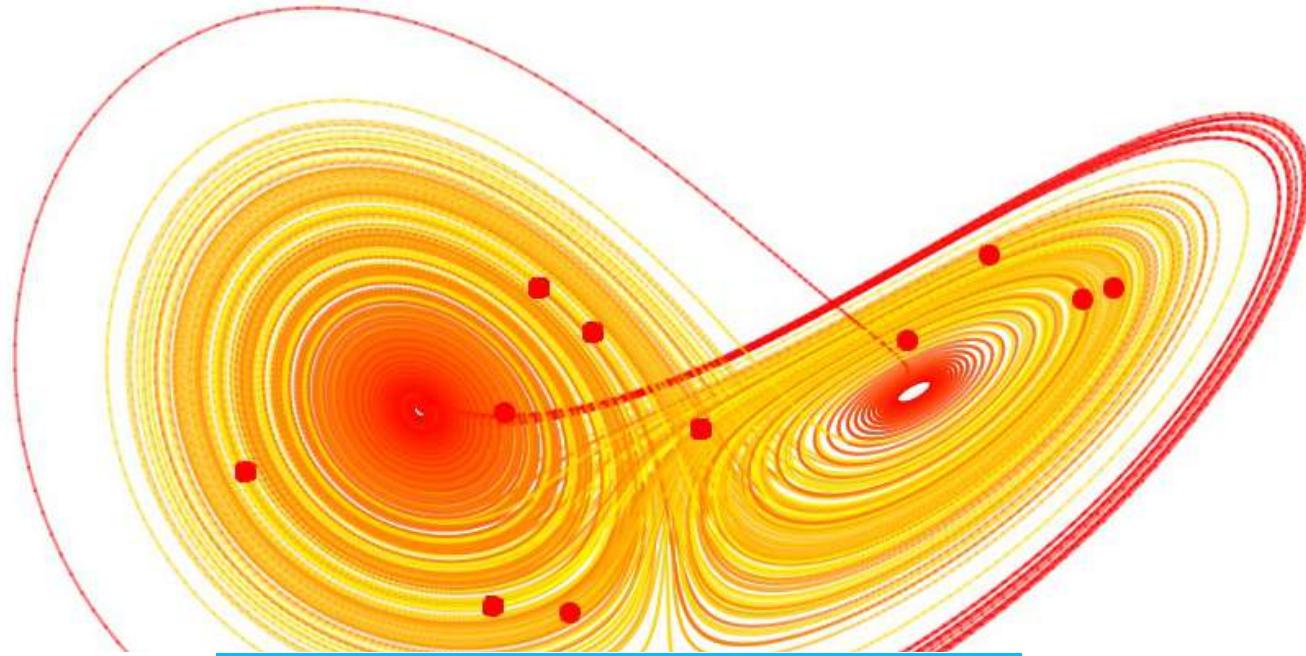
$$\tilde{R}(U_L, x_L, t_L) = -\tilde{M}(U_L)$$

ML with Solver-in-the-loop



# Smarter ML: Learning optimal strategies for dynamical systems from observations

”



**The briefest of brief intros to RL**

# Markov Decision Process (MDP)

- Extends Markov chains **with actions & rewards**
- Discrete-time stochastic **control process** defined by the **tuple of actions, states, rewards and discounts**

$\langle S, A, P(\cdot), R(\cdot), \gamma \rangle$ , with

$a \in A$  : actions

$s \in S$  : state of the system

$\gamma \in \mathcal{R}$  : discount factor

$P(s_{k+1} = s' | \langle s_k = s, a_k = a \rangle)$  : transition probability

# Markov Decision Process (MDP)

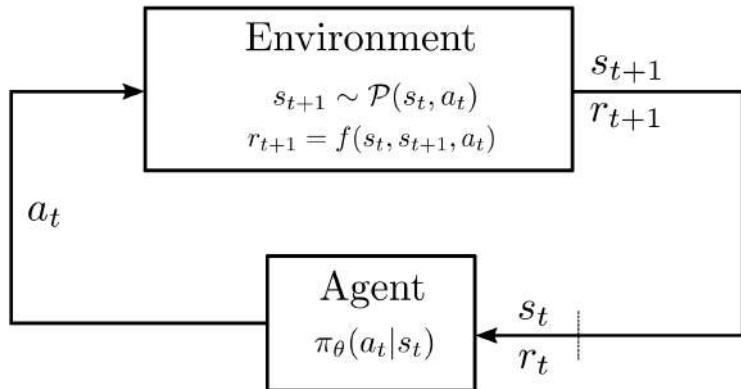
- Markovian IFF memoryless:

$$P(s_{k+1} | \langle s_k, a_k \rangle) = P(s_{k+1} | \langle s_0, a_0 \rangle, \dots, \langle s_k, a_k \rangle), \forall (s_\tau \in S, a_\tau \in A), \tau \{1, \dots, k\}$$

- Formally solved by a **policy**: mapping from state-action space **to the probability of taking action a when in state s**:

$$\pi = S \times A \rightarrow \mathcal{R}, \langle s, a \rangle \rightarrow [0, 1]$$

# Markov Decision Process (MDP)



$$\pi = S \times A \rightarrow \mathcal{R}, \langle s, a \rangle \rightarrow [0, 1]$$

Solved for the optimal policy, e.g. by [reinforcement learning](#)

# Policy Optimization

- Find policy through optimization under uncertainty: **Policy gradient** vs. Q learning
- Some advantages: robustness, stochasticity, continuous action & state spaces
- Maximize the expected **return of an episode**

$$J(\theta) = E_{\theta} \left( \sum_{k=1}^T \gamma^k r_k \right)$$

- **Optimization problem:**

$$\max_{\theta} E_{\theta} \left( \sum_{k=1}^T \gamma^k r_k \right), \text{ subject to } \langle S, A, P(\cdot), R(\cdot) \rangle$$

- **Task:** find  $\theta^*$  such that  $\pi_{\theta^*}(a_t|s_t)$  **maximizes the collected reward**  $r_t$
- Use good old **Gradient Ascent!**

$$\theta^{new} = \theta^{old} + \alpha \nabla_{\theta} J(\theta)$$

- We have to estimate the **steepest gradient with respect to the weights**  $\theta$ :

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \underbrace{\left( \sum_{k=1}^N \gamma^k r_k \right)}_{\text{Cum. reward over } \tau} \underbrace{\nabla_{\theta} \log \pi_{\theta}(\tau)}_{\text{Grad. of the policy}} \right]$$

- Approximate  $\mathbb{E}_{\tau \sim \pi_{\theta}}$  by sampling some **discrete trajectories**  $\tau^{(i)}$ :

$$\tau^{(1)} = \{(s_0, a_0), (s_1, a_1, r_1), \dots, (s_N, a_N, r_N)\}$$

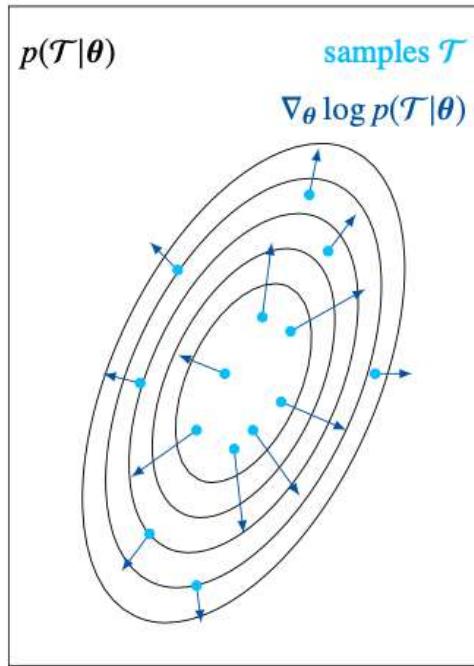
# Policy Optimization

- This is the vanilla policy gradient : VPG
- Improvements: **Trust Region PO**: the policy update is limited so the **KL divergence** remains in a certain bound

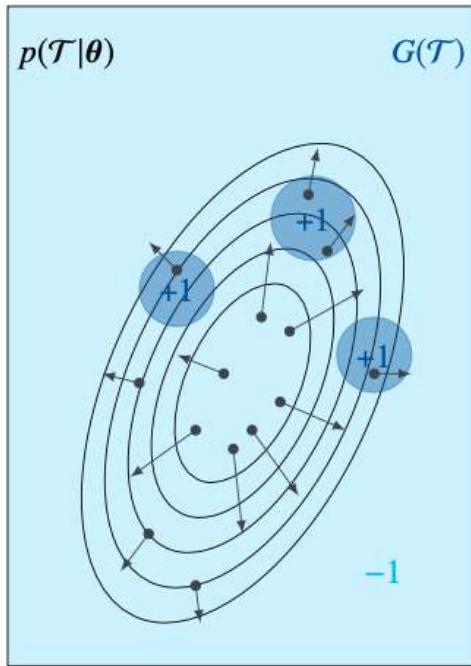
$$\theta' \leftarrow \operatorname{argmax}_{\theta'} \nabla_{\theta} J(\theta)^T (\theta - \theta'),$$

such that  $D_{KL}(\pi_{\theta'}(\cdot|s) || (\pi_{\theta'}(\cdot|s))) \leq \epsilon \forall s \in S$

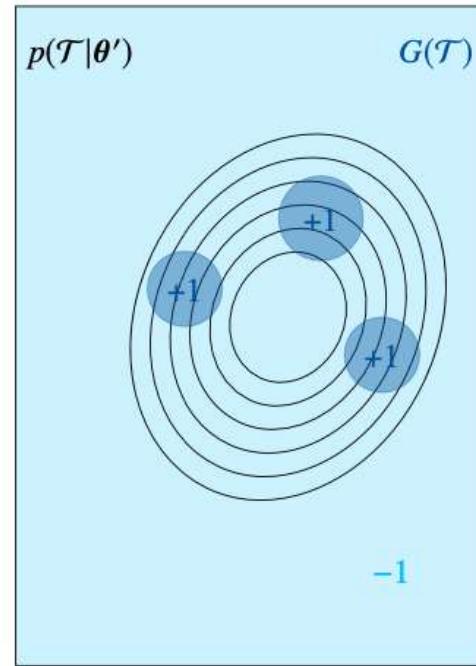
- Adding a baseline to the sampled return: **Advantage function**
- Clipping the objective function: **Proximal Policy Optimization**



**(a) Samples and score functions**

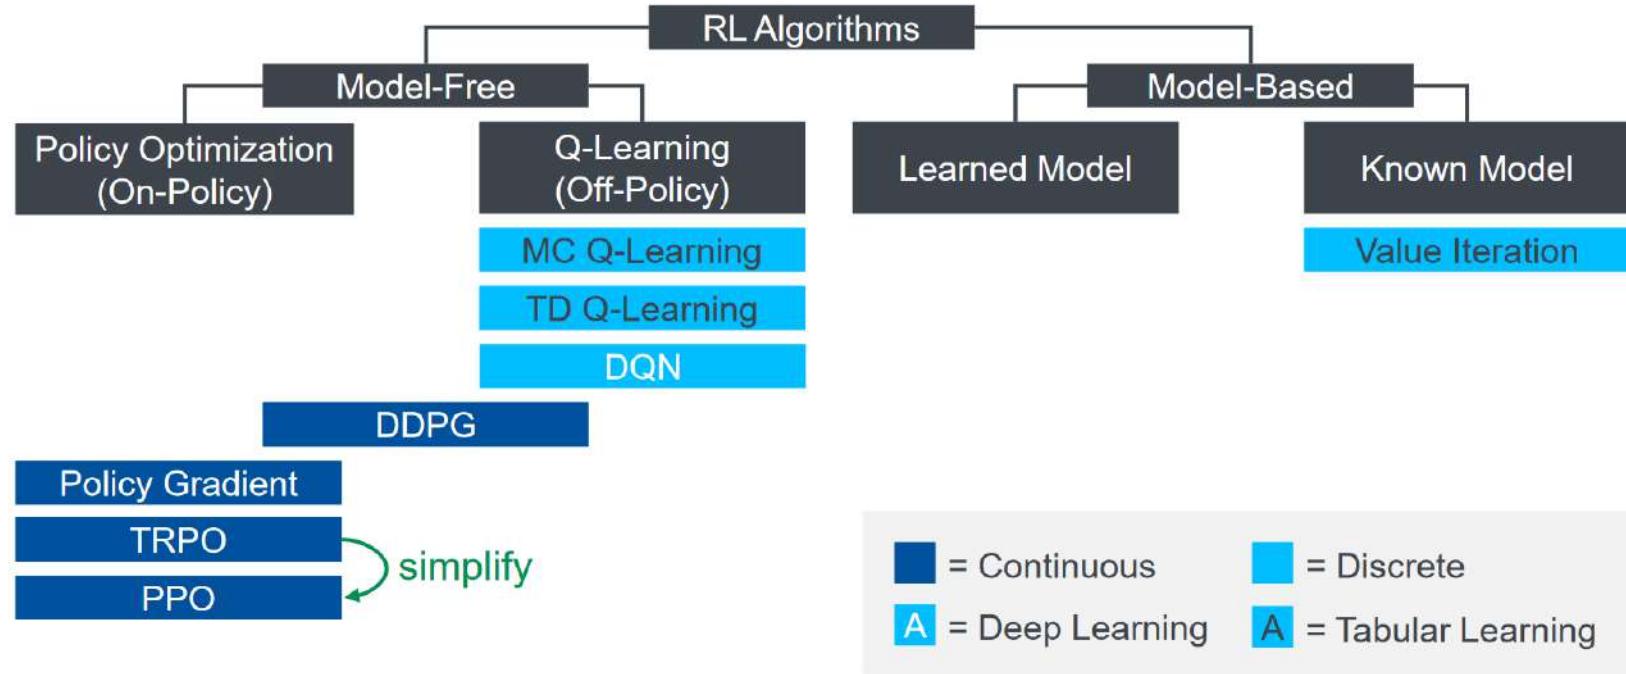


**(b) Return superimposed to Fig. 5a**



**(c) Updated distribution**

With permission from: Hierarchical Reinforcement Learning Approach Towards Autonomous Cross-Country Soaring, AIAA, Stefan Notter, IFR Uni Stuttgart

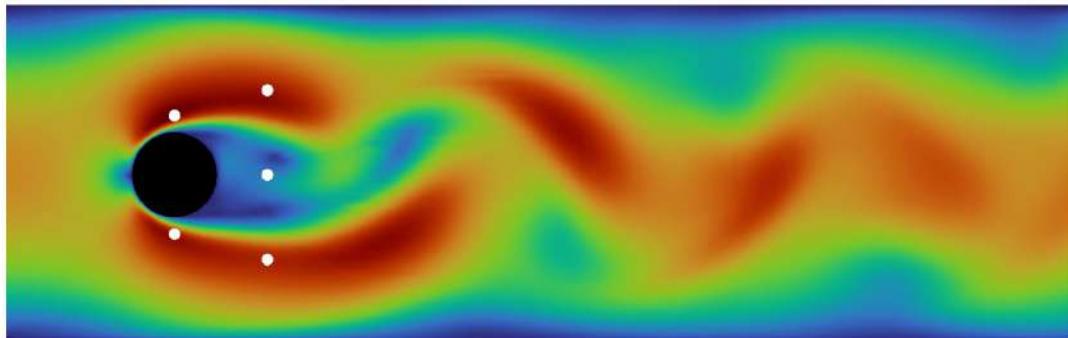
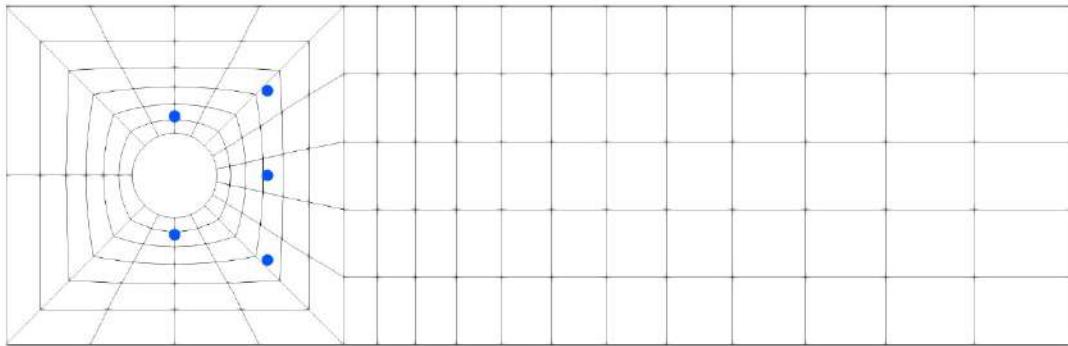


# A simple example first

“ ”

## Example 1: Flow control

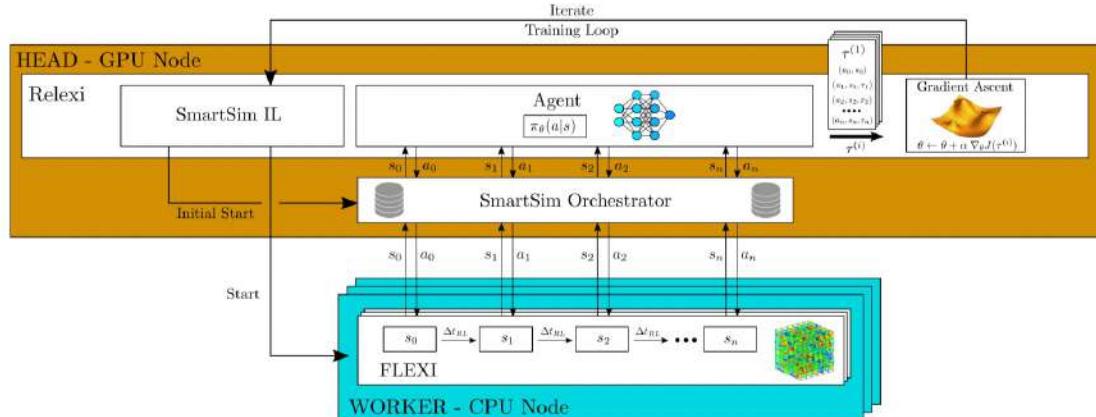
- Tight coupling of “learning” and “simulating”:



- Cylinder at  $Re = 100$ ,  $Ma=0.2$
- 5 Pressure probes as sensors
- Actions: Jets at the poles
- Reward: **Drag reduction**
- Training takes about 2h

# Reinforcement learning framework – ReLeXI<sup>1</sup>: IAG, HLRS & HPE

- Distribution on hybrid HPC systems via the SmartSim Library<sup>2</sup>
- Dedicated GPU node („Head“) for training and model evaluation
- FLEXI instances interactively distributed across multiple CPU nodes („Workers“)
- Communication via in-memory database with the Redis library



<sup>1</sup>Kurz et al., Relexi—A scalable open source reinforcement learning framework for high-performance computing. *Software Impacts*, 2022

<sup>2</sup><https://github.com/CrayLabs/SmartSim>



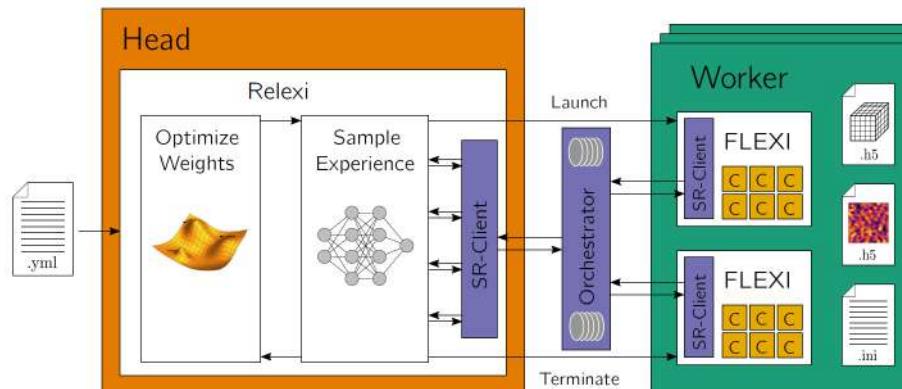
Philipp Offenhäuser &amp; Christopher Williams

# Closing the gap between High-Performance Computing (HPC) and artificial intelligence (AI)

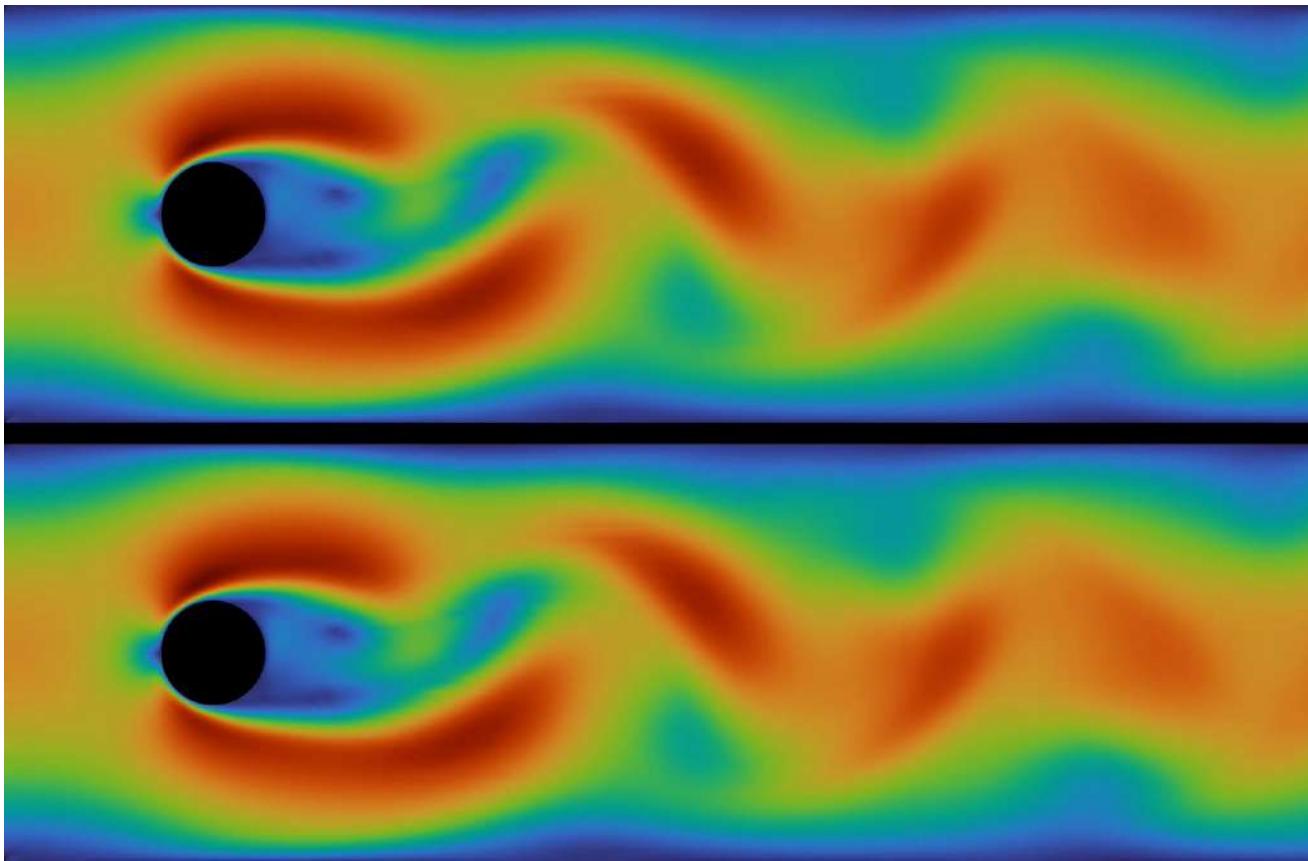
September 15, 2023



Scientists and engineers face a number of technical hurdles to utilizing artificial intelligence (AI) techniques alongside high-performance scientific computing applications. Recently, a multidisciplinary team of researchers at Hewlett Packard Enterprise (HPE), the Institute of Aerodynamics and Gas Dynamics of the University of Stuttgart, and the High-Performance Computing Center Stuttgart (HLRS), created Relexi to overcome these technical challenges and provide the community with a large-scale simulation framework that can be tailored to their use cases. This solution was built on top of another existing open-source library developed at HPE called [SmartSim](#).



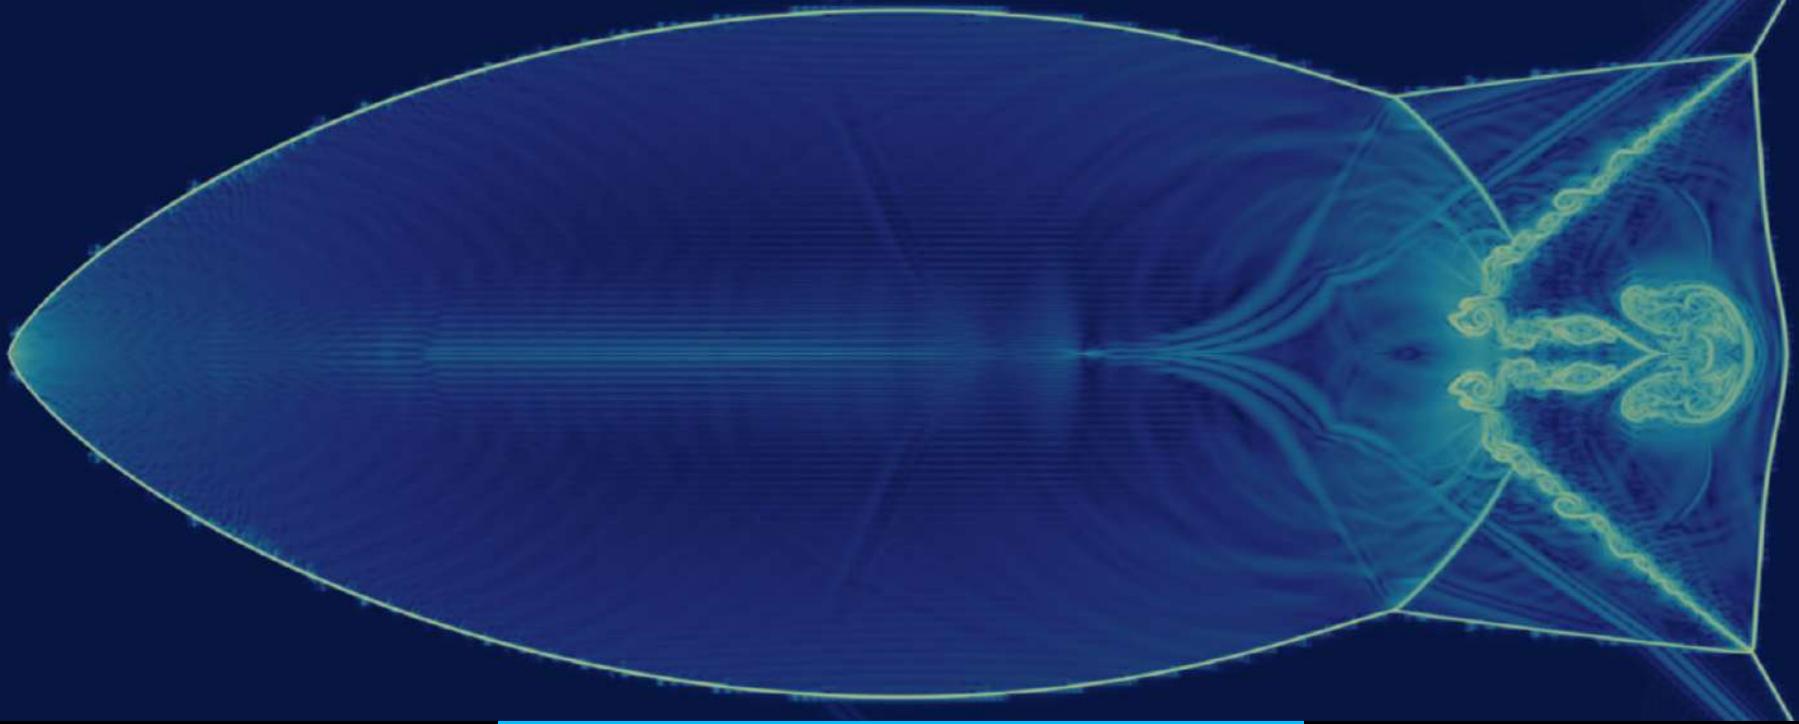
Kurz, Marius, et al. "Deep reinforcement learning for computational fluid dynamics on HPC systems." *Journal of Computational Science* 65 (2022): 101884.



- Jets suppress vortex shedding
- Drag reduction of 10%
- Stability issues with PPO
- Toy problem!

**SL: an offline HPC  
problem.**  
**CFD + RL = an online  
HPC problem.**

“ ”



**RL for discretization consistent LES  
closures**

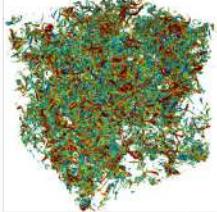
# The problem



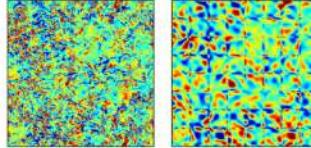
- The LES formulation describes a dynamical, chaotic system
- The computational SGS fluxes make it even harder
- The discrete operators can be **non-linear** for underresolved scales (by design)
- In this context, find discretization-consistent closure strategies through **optimization**
- Back to the Forced HIT
  - ALDM (genetic alg.), Filter-DG (Nelder-Mead), ...

# LES Closures as an MDP

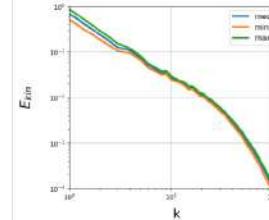
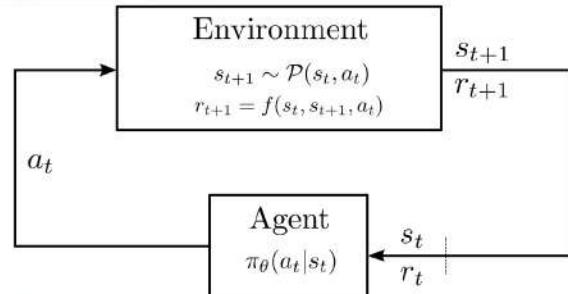
Forced LES of HIT



Implicitly filtered LES with a HO DGSEM



FLEXI iLES



Expected Spectrum



Policy net predicts  $C_s$

Incorporate physics, optimize under discretization

$$\overline{R(u)} = \frac{\partial \bar{u}}{\partial t} + \frac{1}{2} \frac{\delta \bar{u} \bar{u}}{\delta x} = C_1 + C_2 + \frac{1}{2} \frac{\partial}{\partial x} \underbrace{(\widehat{\bar{u} \bar{u}} - \bar{u} \bar{u})}_{C_3^*[\widehat{\cdot}; \cdot]}; \delta_x f = \partial_x \hat{f}$$

- Reynolds stresses: What classical SGS models approximate

- **Complete**, discretization-consistent closure model for full SGS / SFS terms

Optimized eddy viscosity

Optimized dissipative properties

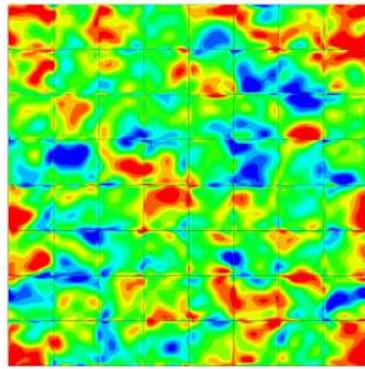
Formulate modelling strategy as an MDP,  
solve with RL

## Things to note

- We find models for the **complete closure terms**  $C_1 \dots C_3$  (and others)
- We **do not need to know**  $C_1 \dots C_3$  a priori: We do not need DNS data, only a statistic
- Instead, “training” means running **a number of LES** simulations
- We optimize the **solution of the overall PDE** (i.e. the LES), not individual terms
- We **discover** new modelling schemes through RL: RL itself is not a model

# Two modelling scenarios

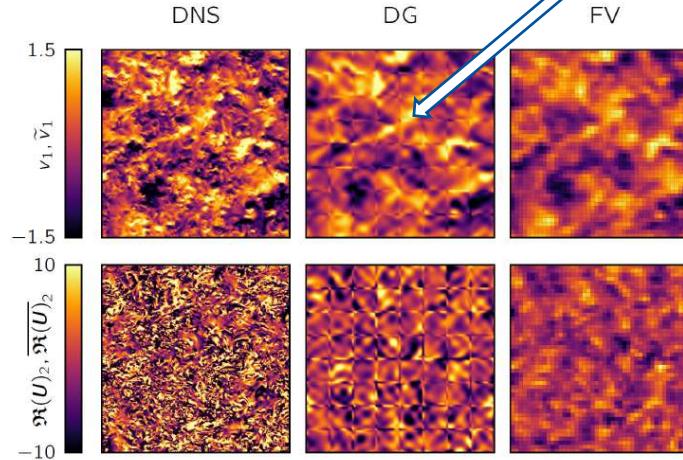
Implicit and **explicit** SGS modeling: Optimal Smagorinsky constant



SSM: For each grid element  $q_{rst}$  with  $r, s, t \in [1, nElems]$  :  $C_s^q = C_s = \text{const.}$

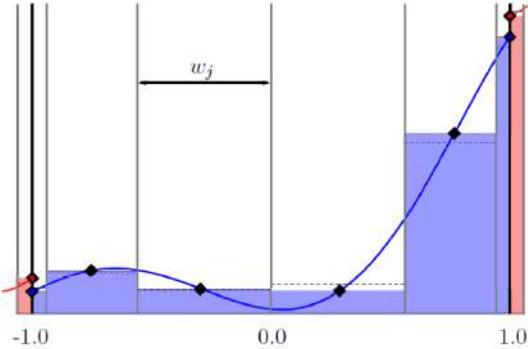
Cs0: For each grid element  $q_{rst}$  with  $r, s, t \in [1, nElems]$  :  $C_s^q = C_s(t)$

Cs2: For each grid element  $q_{rst}$  with  $r, s, t \in [1, nElems]$  :  $C_s^q = \sum_{i,j,k=0}^2 c_i(t) L_i(x) L_j(y) L_k(z)$



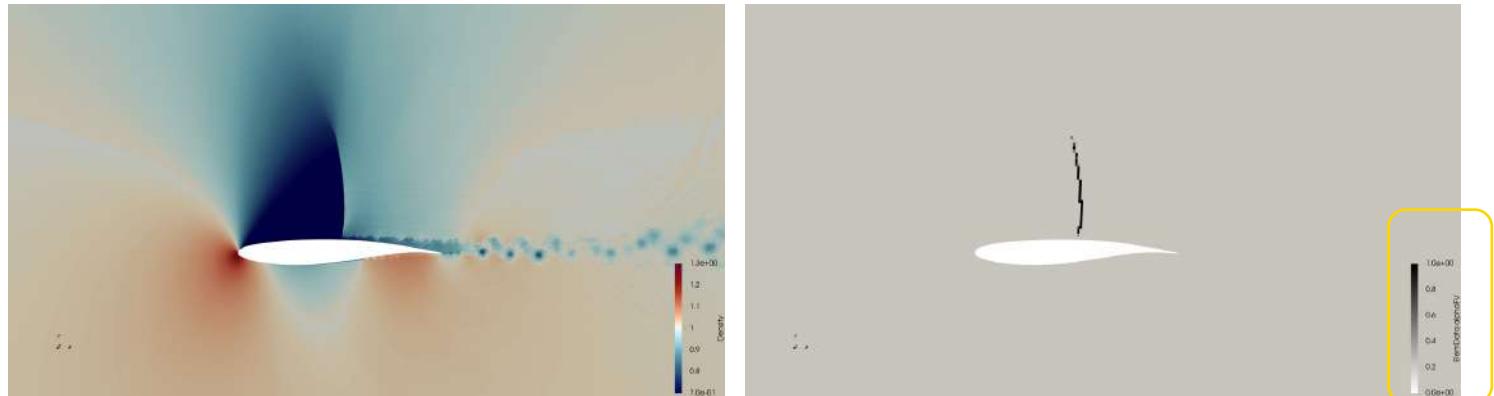
## Two modelling scenarios

Implicit and explicit SGS modeling: Optimal Operator Blending

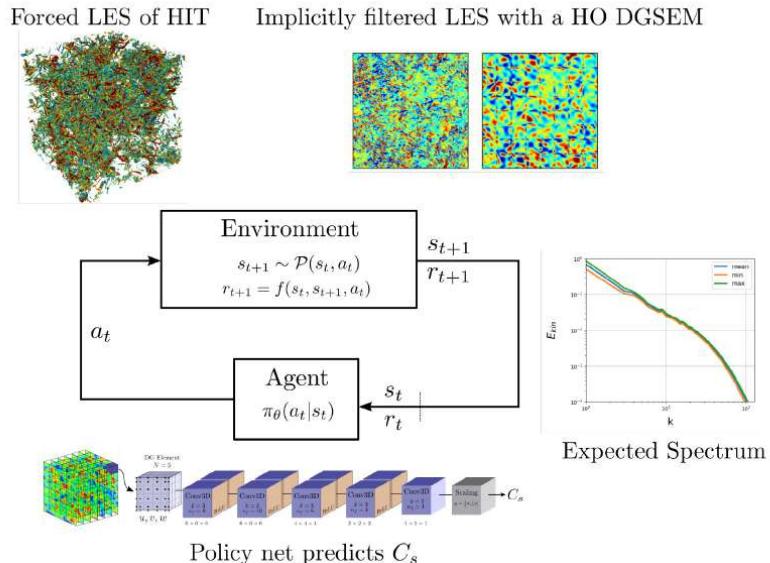


A0: For each grid element  $q_{rst}$  with  $r, s, t \in [1, nElems]$  :  $\alpha_s^q = \alpha(t)$

A2: For each grid element  $q_{rst}$  with  $r, s, t \in [1, nElems]$  :  $\alpha_s^q = \sum_{i,j,k=0}^2 a_i(t) L_i(x) L_j(y) L_k(z)$

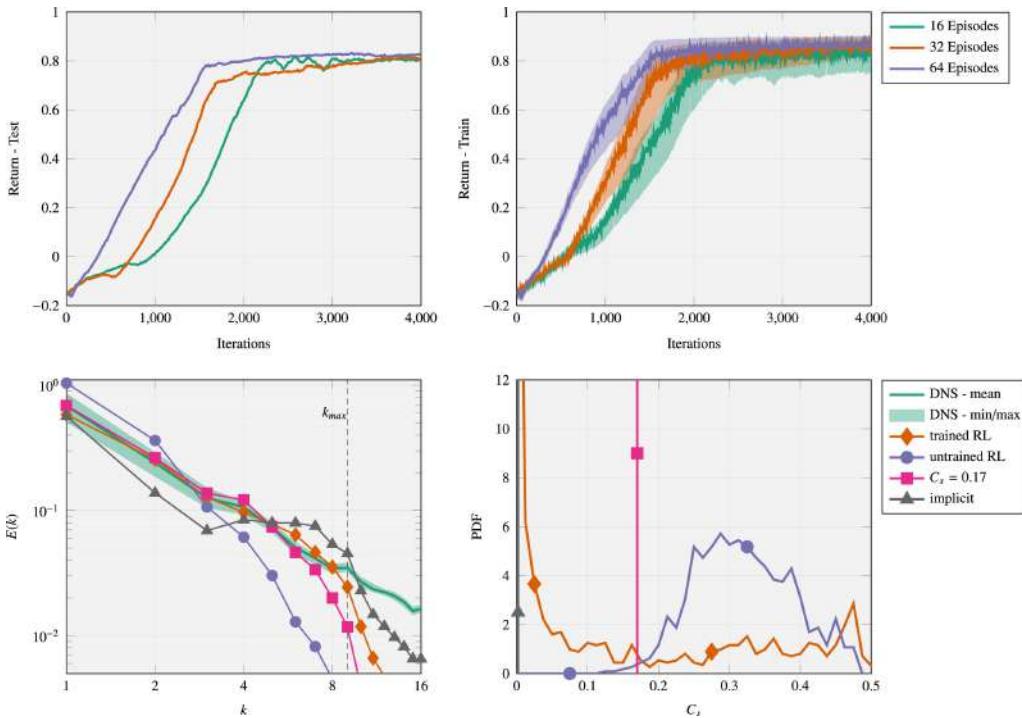


# Optimization Details



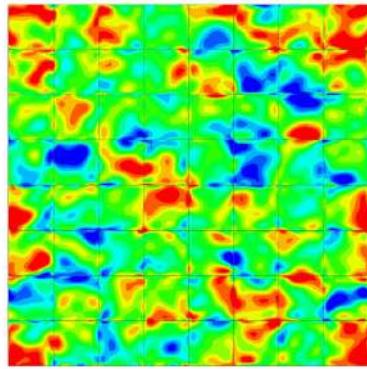
- Environment: Implicitly filtered HO DGSEM
- Reward / Goal: Spectrum up to cut-off
- Agent: Convolutional / Residual NN
- Policy: Choose Cs or Alpha
- No DNS data; optimization requires LES runs only
- Implement as hybrid CPU/GPU parallel code in SmartSim on HAWK-AI
- Training wall clock time: 24h

# Example: Successful optimization for Cs0 case



# Two modelling scenarios

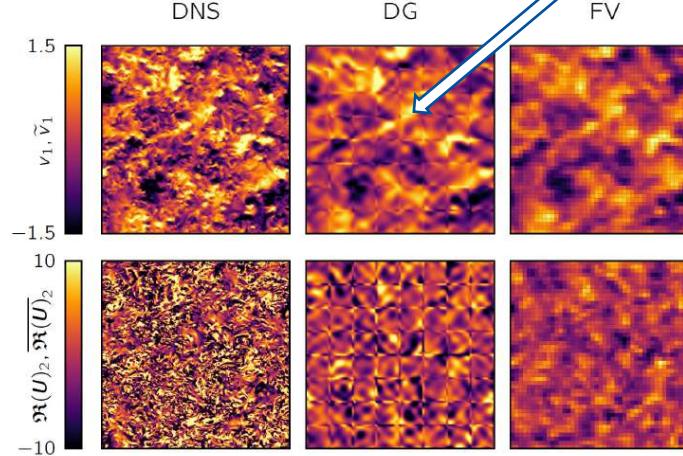
Implicit and **explicit** SGS modeling: Optimal Smagorinsky constant



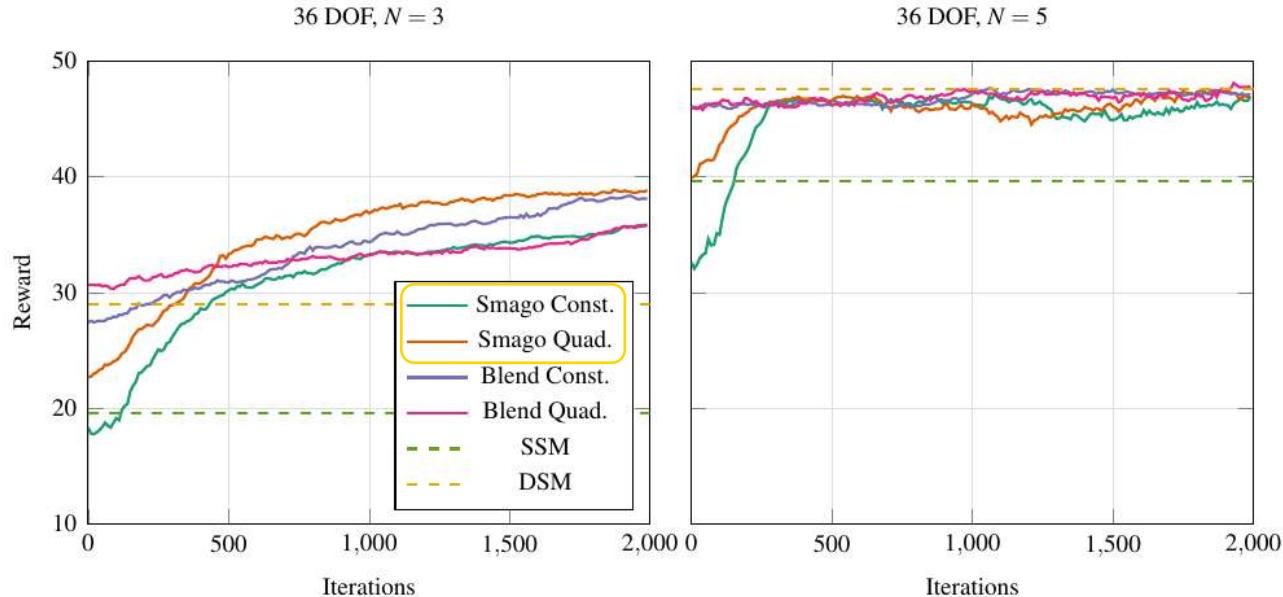
SSM: For each grid element  $q_{rst}$  with  $r, s, t \in [1, nElems]$  :  $C_s^q = C_s = \text{const.}$

Cs0: For each grid element  $q_{rst}$  with  $r, s, t \in [1, nElems]$  :  $C_s^q = C_s(t)$

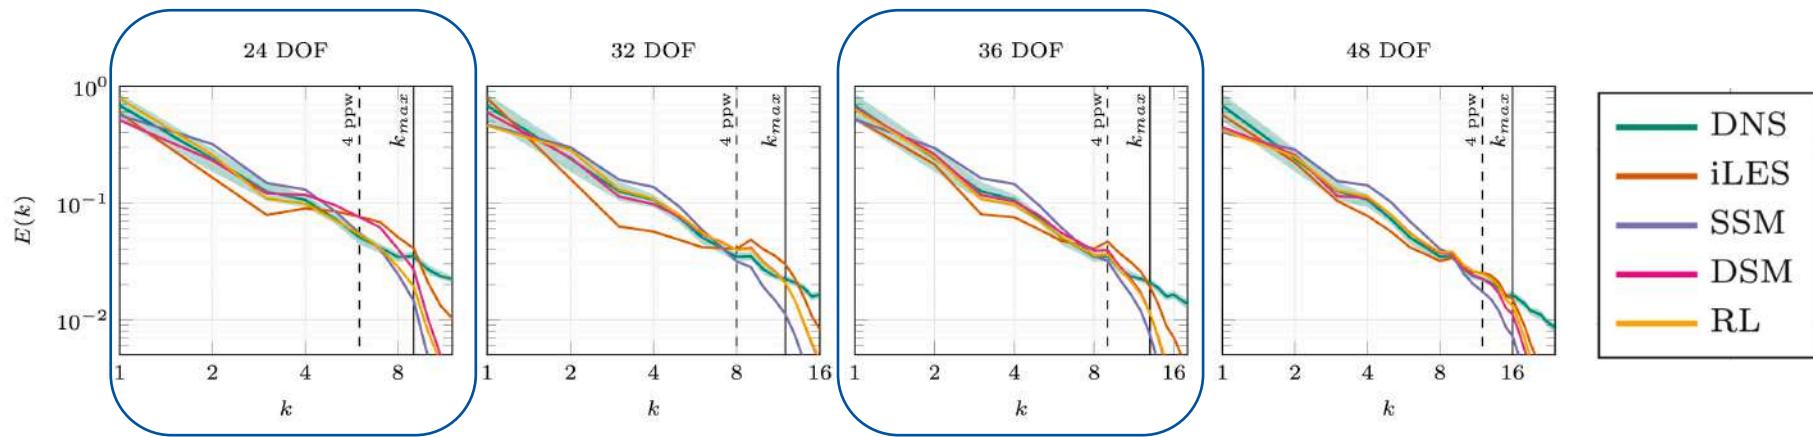
Cs2: For each grid element  $q_{rst}$  with  $r, s, t \in [1, nElems]$  :  $C_s^q = \sum_{i,j,k=0}^2 c_i(t) L_i(x) L_j(y) L_k(z)$



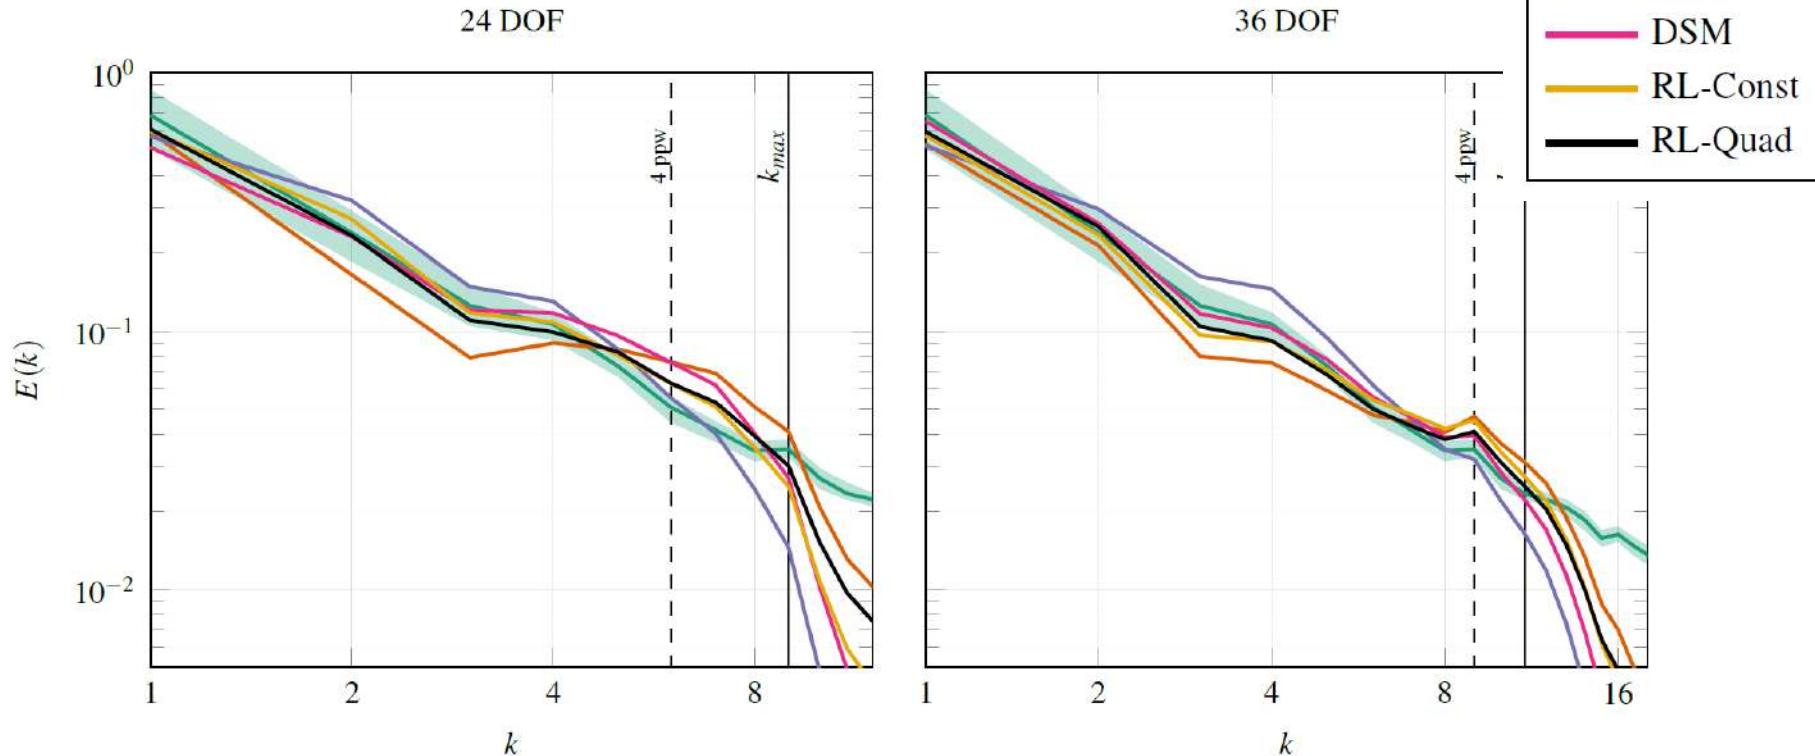
# Results: Explicit Closure with optimized Cs



# Results: Explicit Closure with optimized Cs N=5 for different resolutions

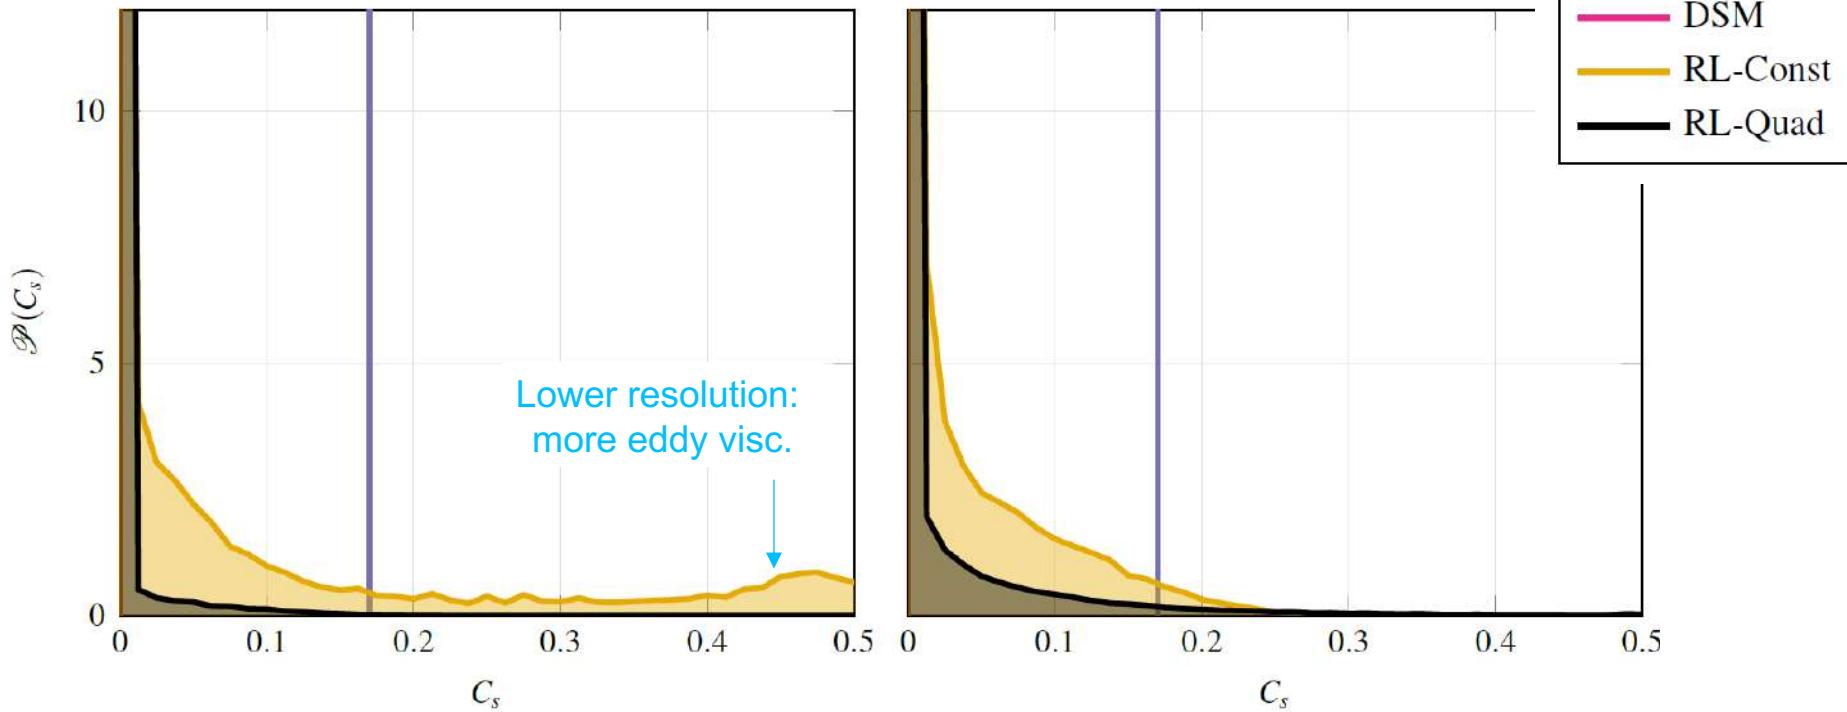


# Results: Explicit Closure with optimized Cs N=5 for different resolutions



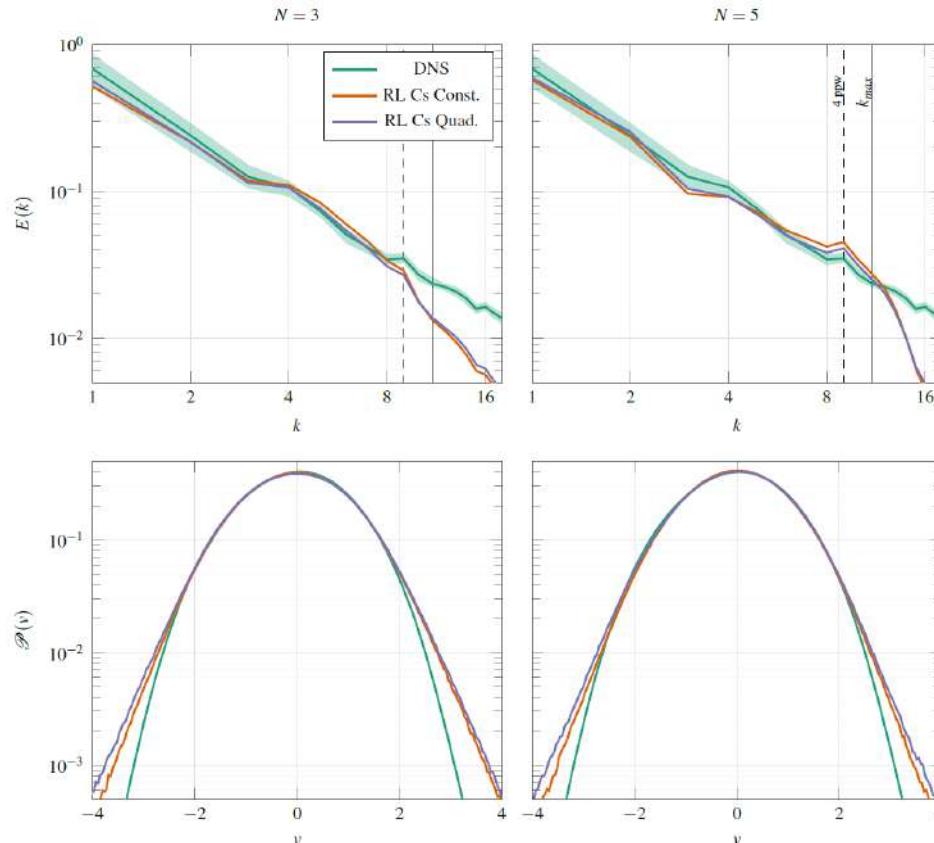
# Results: Explicit Closure with optimized Cs

N=5 for different resolutions



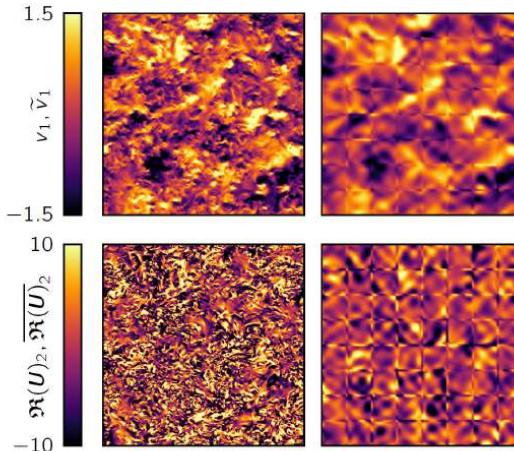
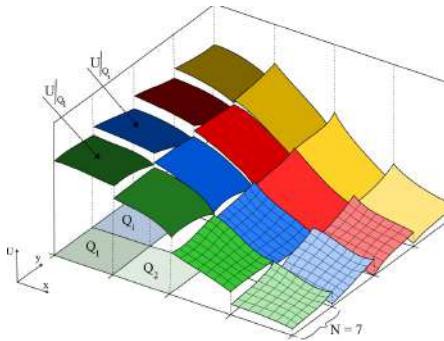
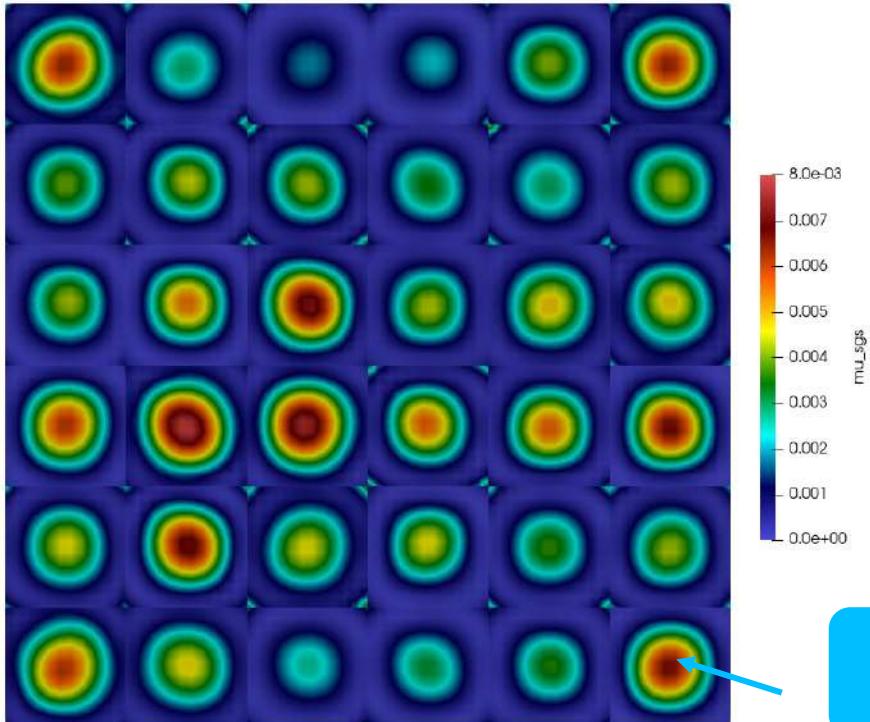
# Results: Explicit Closure with optimized Cs

Different operators, same resolution



# Results: Explicit Closure with optimized Cs

## Mean Eddy viscosity field for Cs2



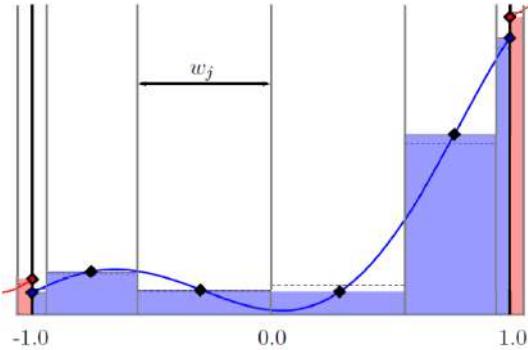
The discretization / induced filter footprint!

A model that accounts  
for the error of the  
discretization.

”

## Two modelling scenarios

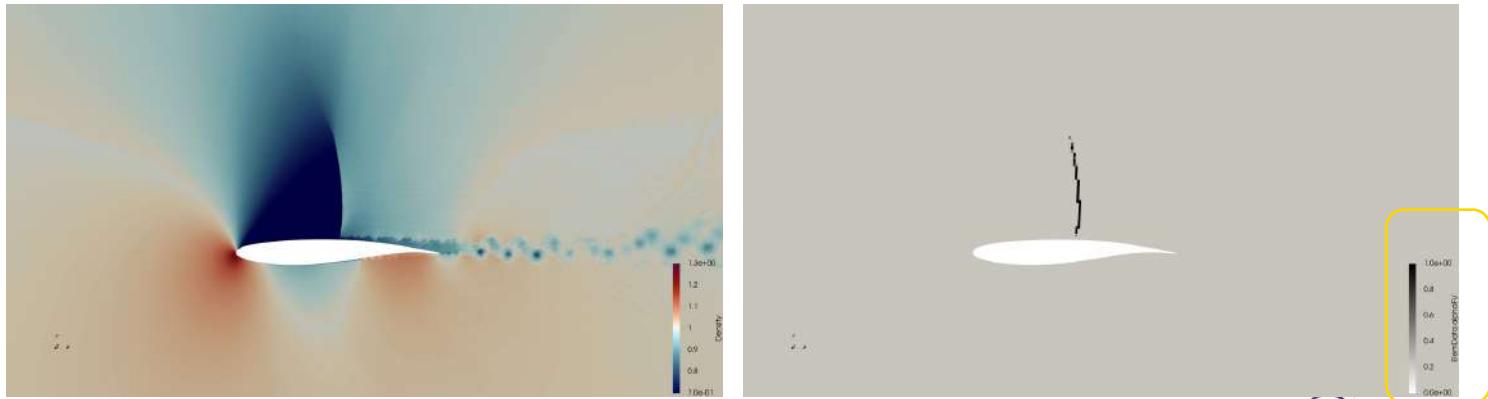
Implicit and explicit SGS modeling: Optimal Operator Blending



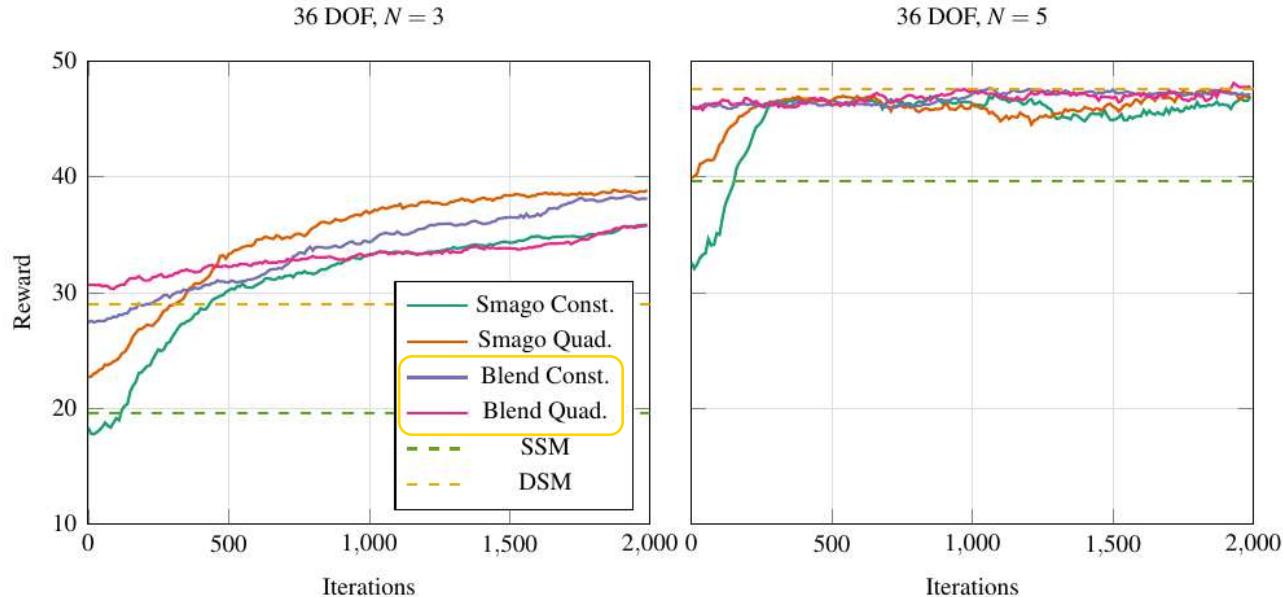
$$\underline{u}^{BL} = \underline{u}^{DG} (1 - \alpha) + \underline{u}^{FV} \alpha$$

A0: For each grid element  $q_{rst}$  with  $r, s, t \in [1, nElems]$  :  $\alpha_s^q = \alpha(t)$

A2: For each grid element  $q_{rst}$  with  $r, s, t \in [1, nElems]$  :  $\alpha_s^q = \sum_{i,j,k=0}^2 a_i(t) L_i(x) L_j(y) L_k(z)$

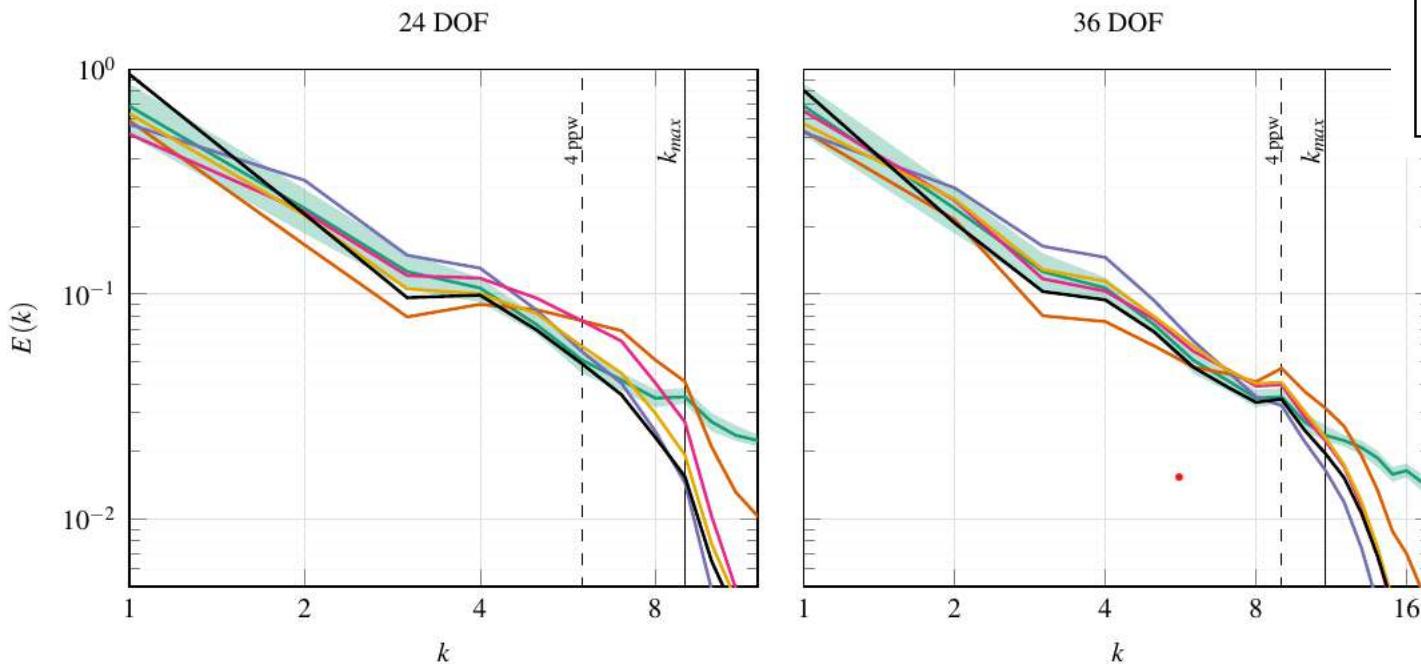


# Results: Implicit Closure with optimized Blending



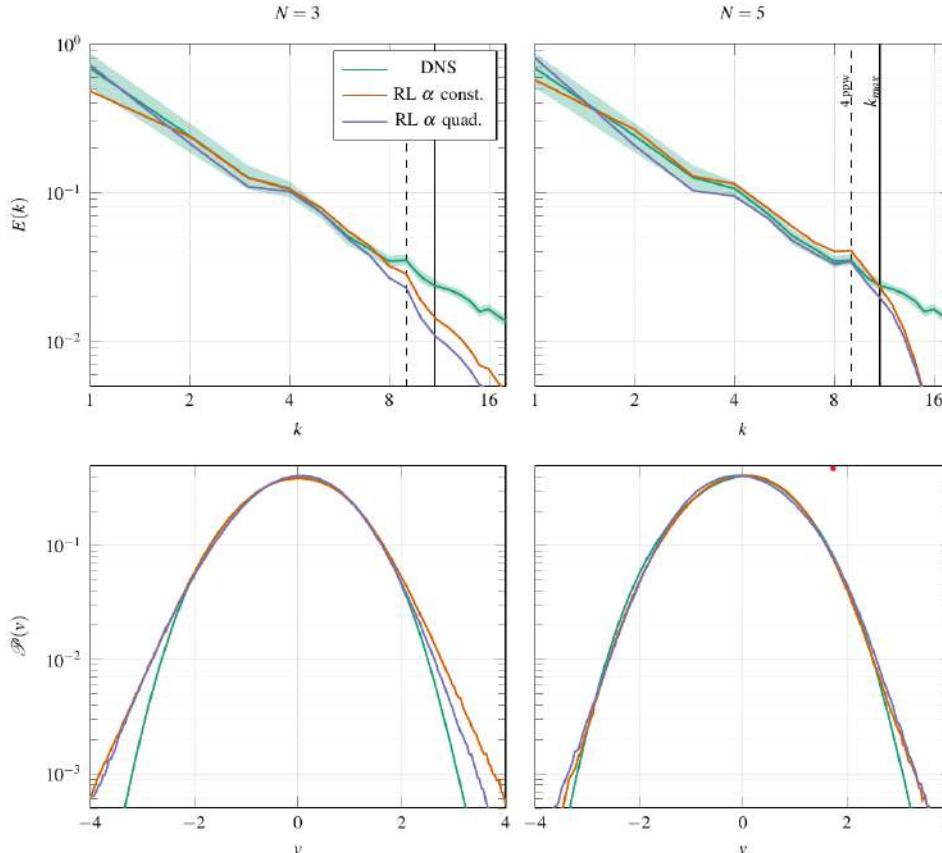
# Results: Implicit Closure with optimized Blending N=5 for different resolutions

- DNS
- iLES
- SSM
- DSM
- RL-Const
- RL-Quad



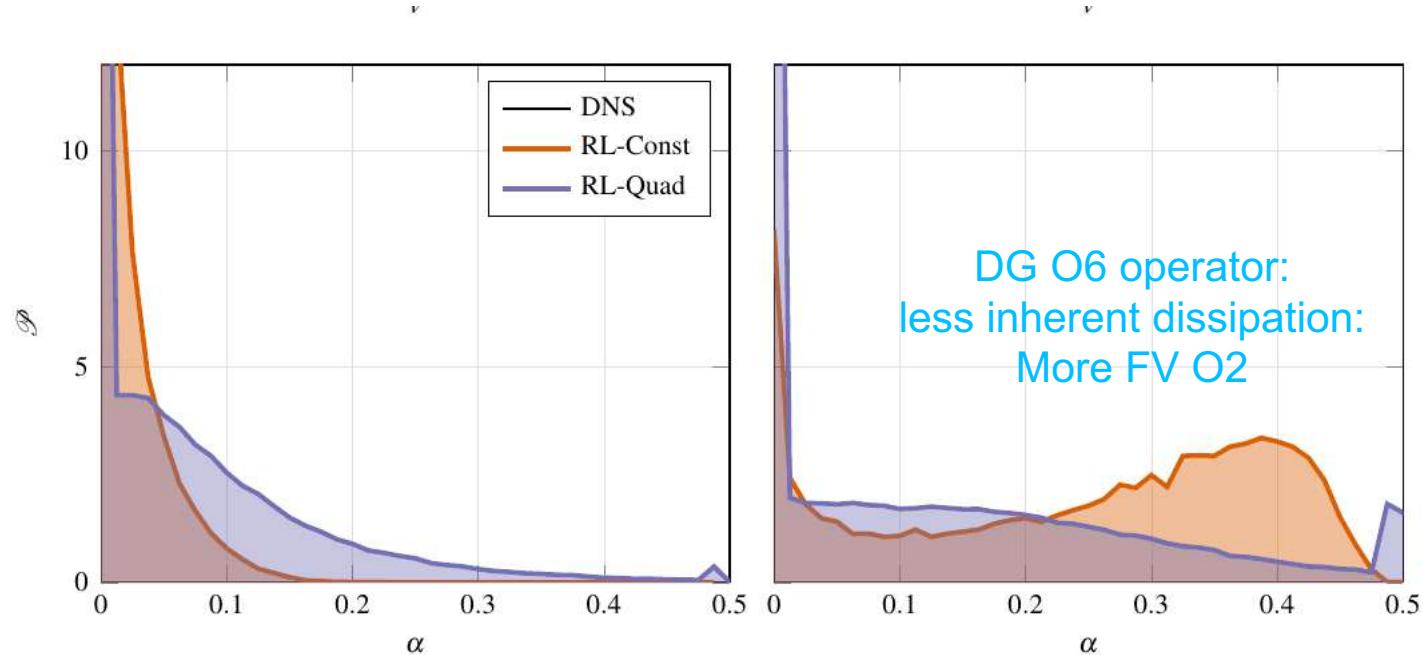
# Results: Implicit Closure with optimized Blending

## Different operators, same resolution

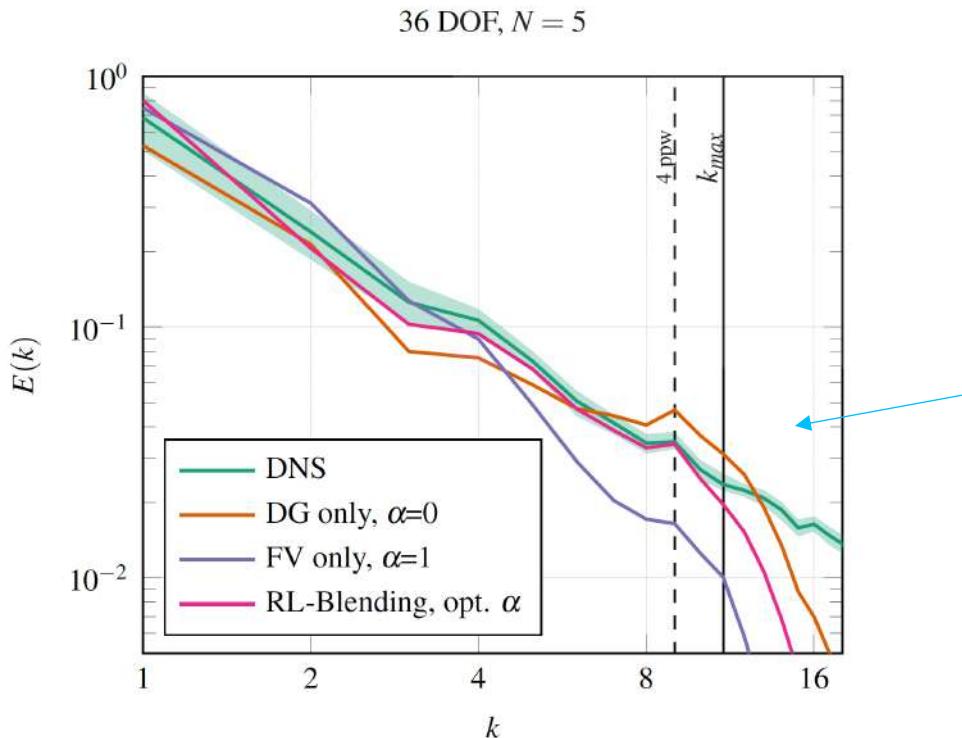


# Results: Implicit Closure with optimized Blending

Different operators, same resolution



# Results: Implicit Closure with optimized Blending



DG MILES: implicit LES model through non-linear combination of HO and TVD scheme

# An “intelligent” LES solver?

”



# Summary

## Food for thought

- Multiscale problems require multi-X numerical schemes
- SL works great if one can define good data a priori
- The classical LES formulation omits the challenges of “real life” LES: C1...C3
- Reinforcement learning with a PDE solver in the loop: Versatile, inclusive optimization environment for flow control and modelling
  - RL does not need fine scale data (like DNS), but statistics. Training happens on the coarse (LES) level: Huge computational savings
  - RL allows to optimize and discover new closures, it is not a model itself: All the knowledge about a “good” model is still applicable
  - RL is an alternative to adjoints or differentiable solvers: It approximates gradients of the cost function through sampling
  - RL is ML: it has hyperparameters and we lack theory in parts

# Our publications on data-driven CFD

- Journal of Computational Physics 398, 108910
- GAMM-Mitteilungen 44
- Journal of Computational Physics 313, 1-12
- Journal of Computational Physics, 423, 109824
- Electronic Transactions on Numerical Analysis, 56, 117–137
- International Journal of Heat and Fluid Flow 99, 109094
- Journal of Theoretical and Computational Acoustics 27, 1850044
- Journal of Computational Physics 441, 110475
- SIAM Journal on Scientific Computing 42 (4), B1067–B1091
- Journal of Computational Science 65, 101884
- Journal of Computational Physics 408, 109303
- Software Impacts 14, 100422
- Computers & Fluids 228, 105039
- Wear 508, 204476
- Journal of Open Source Software 8 (82), 4683
- PAMM 23 (1), e202200207
- Toward Discretization-Consistent Closure Schemes for Large Eddy Simulation Using Reinforcement Learning  
[arXiv:2309.06260v1](https://arxiv.org/abs/2309.06260v1)

# What do we need LES for?

First DNS of self-similar, compressible turbulent boundary layer with pressure gradients at sub- and supersonic Mach numbers.



**University of Stuttgart**  
Germany

Simulation is performed at the *Institute of Aerodynamics and Gas Dynamics* (IAG) of the University of Stuttgart by:  
Wenzel, Gibis, Kloker and Rist.



# Thank you!



**Andrea Beck**

e-mail [beck@iag.uni-stuttgart.de](mailto:beck@iag.uni-stuttgart.de)  
phone +49 711 685-60218  
[www.iag.uni-stuttgart.de](http://www.iag.uni-stuttgart.de)  
[numericsresearchgroup.org](http://numericsresearchgroup.org)

- Kurz, Marius, Philipp Offenhäuser, and Andrea Beck. "Deep reinforcement learning for turbulence modeling in large eddy simulations." *International Journal of Heat and Fluid Flow* 99 (2023): 109094.
- Kurz, Marius, et al. "Relexi—A scalable open source reinforcement learning framework for high-performance computing." *Software Impacts* 14 (2022): 100422.
- Mossier, Pascal, Andrea Beck, and Claus-Dieter Munz. "A p-adaptive discontinuous Galerkin method with hp-shock capturing." *Journal of Scientific Computing* 91.1 (2022): 4.
- Schwarz, Anna, et al. "A Reinforcement Learning Based Slope Limiter for Second-Order Finite-Volume Schemes." *Proceedings in Applied Mathematics and Mechanics* 7 (2022): 10.
- Kopper, Patrick, et al. "Boundary-Layer Dynamics in Wall-Resolved LES Across Multiple Turbine Stages." *AIAA Journal* 59.12 (2021): 5225-5237.
- Dürrwächter, Jakob, et al. "A high-order stochastic Galerkin code for the compressible Euler and Navier-Stokes equations." *Computers & Fluids* 228 (2021): 105039.
- Beck, Andrea, and Marius Kurz. "A perspective on machine learning methods in turbulence modeling." *GAMM-Mitteilungen* 44.1 (2021): e202100002.